

Linguistic transparency as mediated accessibility

Brett Reynolds *

Humber Polytechnic & University of Toronto

Jong-Bok Kim

Kyung Hee University

8th May 2026

Abstract

Linguistic transparency is best understood as a family of mediated accessibility relations rather than a unitary syntactic property. A construction, expression, or grammatical configuration is transparent when information that might otherwise be encapsulated remains accessible for agreement, interpretation, constructional generalisation, processing, or cross-linguistic comparison. The umbrella schema *X is transparent w.r.t. property P for relation R when P remains accessible through X for the purposes of R* is a comparative concept in Haspelmath's sense: an analyst's question, applicable to any language, that asks which configurations let an internal property reach a grammatical or interpretive relation.

The schema's job is to locate PROJECTIBILITY PROFILES: cases where an observable base licenses non-trivial projection to a target, stabilised by some identifiable source. Different transparency profiles are stabilised differently: lexical convention, constructional inheritance, usage-based entrenchment, processing architecture, formal-semantic environment. The paper's central question isn't whether transparency is a kind but when transparency attributions support induction and what stabilises that induction. CGEL's number-transparent quantificational nouns are the strongest profile worked through: a small, relatively stable class whose membership predicts agreement override, characteristic determiner patterns, partitive-vs-non-partitive *of*-complement availability, fused-head behaviour, count/non-count percolation, and premodification tendencies, stabilised by conventional lexical inheritance. The form-meaning accessibility family across constructional compositionality, compound rated transparency, and morphological priming licenses within-domain prediction in three operational windows; whether the family constitutes a cross-domain projectibility cluster on shared items is an open research question, not a demonstrated result. The projectibility profile is the main concept; Slater's stable-property-cluster (SPC) framework is the general diagnostic; Khalidi's causal-network account is reserved as a possible refinement for profiles with demonstrated causal ordering, applied here only where the evidence motivates it. Coordination is excluded as a non-instance; referential transparency is acknowledged as the schema's boundary case.

The paper extends to a relational concept the projectibility-first analysis that related work applies to ontological ones, with kindhood as a downstream verdict on profiles that earn it rather than the master question of the analysis.

*Contact: brett.reynolds@humber.ca

I INTRODUCTION

Two simple sentences make the puzzle visible. *A lot of tourists are in town.* *A lot of money is being spent.* The verb agrees with the *of*-complement (*tourists*, *money*) rather than with the surface head *lot*, which is morphologically singular in both cases. The complement is plural count in the first sentence, singular non-count in the second. The construction lets the relevant feature of the complement reach the verb.

NUMBER-TRANSPARENT QUANTIFICATIONAL NOUNS is the term used in Huddleston and Pullum's *Cambridge Grammar of the English Language* (2002) (henceforth CGEL). Core members include *lot*, *plenty*, *lots*, *bags*, *heaps*, *loads*, *oodles*, *stacks*, *remainder*, *rest*, *number*, and *couple*, with *majority* and *minority* at the boundary. Membership is small and relatively stable, and members head NPs whose agreement override is, on CGEL's analysis, the default pattern, contrasting with the optional override of ordinary collective nouns. Many other configurations behave similarly: a form mediates between an internal property and a grammatical or interpretive relation, and the property remains accessible through the mediator. The umbrella schema for the family follows.

X is *transparent* with respect to property *P* for relation *R* when *P* remains accessible through *X* for the purposes of *R*.

The schema is a tool for a particular question: when is calling something *transparent* doing analytical work, and when is it just a label? Sometimes the label genuinely earns its keep; sometimes it doesn't. The *a lot of* construction shows what earning the label looks like.

The schema's substantive content is the *remains accessible through X* clause. A case fits only if *X* has internal structure whose default behaviour would have made *P* inaccessible to *R*. Cases without that default-conceal background are excluded. Coordination (§2.3) is the textbook exclusion: *Bob and Jane are happy* has no head whose morphological number competes for agreement with the conjuncts' features, so there's nothing the construction could have hidden, and calling coordination *transparent* would be a relabel rather than an analysis. Everyday senses of TRANSPARENT PROSE or TRANSPARENT INTENTIONS, where no specific mediator-property-relation triple is in play, are excluded for a related reason: the schema's slots have no values to bind. The exclusion mechanism is what keeps the schema from cataloguing every use of *transparent* in linguistics, and the rest of the paper is occupied with cases where the slots have non-trivial values.

Once we recognise that *lot*, *plenty*, *lots*, *bags*, *heaps*, *loads*, *oodles*, *stacks*, *remainder*, *rest*, *number*, and *couple* (in its quantificational sense) form a small, conventionally stable class, we can predict more than the agreement override that motivated the label. Members of the class show characteristic patterns of determiner co-occurrence; they admit partitive *of*-complements in some subgroups but not others; they behave distinctively under the fused-head construction; the count-or-non-count profile of the *of*-complement determines the profile of the whole NP; modification by ordinary adjectives is restricted in characteristic ways. Knowing that a noun belongs to the class lets us predict much of the rest, with the predictions reliable for the core members and graded at the boundary (*majority*, *minority*, sense-split *couple*). The label is doing work: a small lexical inventory whose membership predicts a cluster of correlated behaviours, attenuated at the edges rather than guaranteed everywhere.

Cases like this we'll call PROJECTIBILITY PROFILES. A profile has three components: an OBSERVABLE INPUT (here, lexical-class membership); a set of *predictions* the input licenses (here, the cluster of agreement, distributional, and modificational behaviours); and a STABILISING PROCESS that holds

the input together (here, ordinary lexical inheritance). The number-transparent class is a strong example. The predictions are non-trivial; the inventory is stable; the stabilising process is the same one speakers use for any closed grammatical class.

Not every transparency label names a profile this strong. Some name profiles whose predictions are weaker, or whose stabilising process is different, or both. The paper asks which transparency labels do projection-supporting work, and what makes them able to do it.

X (the mediator) can be a lexical item, a partly or fully schematic construction, an intensional context, or a derivationally complex word. P can be a number/count feature, a constituent's meaning, a referent, or a stem's lexical representation. R can be agreement, interpretation, substitution under co-reference, or masked priming. Accessibility itself admits two readings: class-like (core members behave categorically, even where class boundaries are fuzzy) and gradient (compositionality scales, transparency ratings, priming magnitudes).

The argument is constructive. The paper isn't a survey of every use of TRANSPARENCY in linguistics, nor a paper on extraction, islands, barriers, or phases; extraction-transparency is acknowledged in §2 but not developed. It isn't a paper on phonological or orthographic transparency, nor on the everyday senses of *transparent explanation* or *transparent terminology*. Referential or DE RE/DE DICTO transparency receives a controlled mention in §4 as the boundary case where the schema's mediator-form structure stretches thinnest.

MEDIATED ACCESSIBILITY is more specific than CLARITY and broader than COMPOSITIONALITY. Clarity is rhetorical: it doesn't single out a mediator or a relation. Compositionality is the limiting case of FORM-MEANING transparency, where a mediator's parts predict the whole. The umbrella schema is broader: it covers form-meaning cases (constructions, compounds, derived words) and accessibility cases that aren't reducible to whether parts predict wholes. Number transparency is the paradigm of the second type: *lot*'s morphological singular doesn't predict the verb's number, and under the construction it doesn't have to.

The paper unfolds as follows. Section 2 treats morphosyntactic transparency, with CGEL number-transparent quantificational nouns as the paper's strongest projectibility profile (§2.4); collective NPs as a transitional case; coordination as a non-instance; Korean reserved for the co-authored extension. Sections 3 to 5 work through the form-meaning accessibility family: theoretical compositionality (Goldberg, 2006), rated compound transparency (Bell & Schäfer, 2016), and time-sensitive lexical processing (Heyer & Kornishova, 2018). Each licenses within-domain projection; whether they jointly license a cross-domain projectibility cluster is an open research programme, not a demonstrated result. Section 4 also handles referential transparency as the schema's boundary case. Section 6 invokes Haspelmath's (2010) descriptive-versus-comparative-concept distinction to keep cross-linguistic remarks honest. Section 7 names the stabilising structures behind each profile and uses Slater's (2015) and Khalidi's (2018) frameworks from the philosophy of natural kinds as diagnostic tools for evaluating how strongly each profile supports inductive inference.

2 MORPHOSYNTACTIC TRANSPARENCY

Morphosyntactic transparency is the case where the mediator X is a syntactic construction (or a lexical class within a construction) and the relation R is a matching relation between a non-head property

and a co-occurring exponent.¹ Subject-verb agreement is the textbook instance. CGEL’s number-transparent quantificational nouns provide the paradigm in §2.1; collective NPs the transitional case in §2.2; coordination is a non-instance in §2.3. §2.4 argues that the paradigm is a strong candidate projectibility profile, with kindhood as a downstream verdict. §2.5 adds transparent free relatives as a second profile in the same morphosyntactic family, stabilised by constructional rather than lexical inheritance and bringing two dissociable *R*-relations under one mediator. §2.6 [to be completed].

2.1 NUMBER-TRANSPARENT QUANTIFICATIONAL NOUNS

The paradigm again: *A lot of tourists are in town.* *A lot of money is being spent.* The verb agrees with the *of*-complement (*tourists* in the first, *money* in the second) rather than with the surface head *lot*, which is morphologically singular in both. Schema-wise, *X* is the *number-transparent quantificational noun* + *of*-complement construction; *P* is the number/count profile of the *of*-complement; *R* is subject-verb agreement.

CGEL’s “simple agreement rule” is head-driven: a singular head NP triggers singular agreement, a plural head NP plural agreement (Huddleston & Pullum, 2002, §18.1). Number-transparent NPs override the rule. CGEL describes the override as obligatory in the central cases; rare singular-verb attestations with *number* are noted in footnote 73 to §18.2 and analysed there as hypercorrections rather than as variant grammar. The contrast with the optional override of collective nouns (§2.2) is what we’ll take as the operative point. CGEL’s example [9] illustrates with both directions of override: *A number of spots have/*has appeared* (plural override on a singular head); *Heaps of money has/*have been spent* (singular override on a plural head). The construction’s job is to let the *of*-complement’s number reach agreement past the head’s morphological number, which would otherwise have determined it.

The core inventory is small and well documented. CGEL identifies the principal members in example [57]: *lot*, *plenty* (singular forms; *lot* most often with *a*-determiner, *plenty* typically with none); *lots*, *bags*, *heaps*, *loads*, *oodles*, *stacks* (plural forms, typically with no determiner); *remainder*, *rest* (singular forms typically with *the*-determiner); *number*, *couple*₂ (singular forms typically with *a*-determiner, taking only plural *of*-complements, partitive or non-partitive). Of these, “the clear cases of number-transparent singular nouns are *lot*, *number*, and *couple*₂”; *majority* and *minority* are borderline (CGEL [13] notes plural override is “surely obligatory” in some examples but singular agreement is occasionally attested) (Huddleston & Pullum, 2002, §18.2). The determiner patterns above are typical, not absolute: *the lot of it* (in a totality sense) is grammatical alongside *a lot of it*, and similar sense-conditioned variations exist for other members.

Class membership predicts a cluster of behaviours beyond the agreement override. The relevant contrast is between PARTITIVE *of*-complements, which require a definite complement that picks out a subset of a definite set (*of the money*, *of those books*, *of these soldiers*), and NON-PARTITIVE *of*-

¹Extraction transparency — whether an island or barrier is opaque or transparent to long-distance dependency formation — is a morphosyntactic-accessibility relation under the schema (with *R* = dependency formation), and the associated question remains active. Transformational accounts posit empty elements (traces, copies) that mediate the connection between a fronted element and its base position; non-transformational dependency accounts (da Cunha & Gibson, 2025; Gibson, 2025) argue that no such mediating empty elements are needed and that direct head-to-dependent connections better fit acceptability and processing data. Both sides make schema-level claims about whether *X* contains invisible mediating structure. Extraction transparency is acknowledged here as a further case the schema covers but not pursued substantively, since the centre of the present argument is the agreement-style cases.

complements, which take indefinite complements (*of money, of soldiers*). *Lot* most often takes *a* as determiner with limited premodification (*a whole/huge lot of time*) (Huddleston & Pullum, 2002, §3.3); *the lot of it* is also available in a totality sense, but other determiners are blocked (**some lot of it*). *Plenty* typically takes no determiner or modifier (*plenty of butter*, not generally *a remarkable plenty of butter*), and like the plural-only forms it essentially doesn't allow partitive *of*-complements.² The plural-only forms (*lots, bags, heaps, loads, oodles, stacks*) show the same pattern: *They have oodles of money* is normal; *Oodles of the money had already been spent* is rare to non-attested.

Remainder and *rest* take only partitive *of*-complements in their characteristic use: *the remainder of the time*, not *the remainder of time*. *Number* and *couple*₂ take plural obliques: *a number of the protesters*, not *a number of money*. The pattern is highly structured even where individual cases admit sense-conditioned variation.

This shows what number transparency *is* in the umbrella schema's vocabulary. *X* (the construction) doesn't conceal *P* (the number/count profile of the *of*-complement) from *R* (agreement). What *X* does conceal from agreement is the morphological number of the head noun *lot, plenty, heaps*, etc. The agreement-targeting feature is on the *of*-complement, not on the surface head. That's what makes *X* transparent with respect to *P*: *P* is on the complement, not on the head, and *X* lets *P* reach the verb anyway.

2.2 COLLECTIVE NPs AS A TRANSITIONAL CASE

Singular collective nouns (*committee, jury, family, team, government*; also *clergy, intelligentsia, public*, etc.) allow optional plural override of head-driven agreement: *The committee is/are divided* (CGEL [8]) (Huddleston & Pullum, 2002, §18.2). The override is meaning-sensitive: singular agreement focuses on the unit, plural on the members. Plural is unavailable with predicates that apply to the whole but not individual members (*The committee consists/*consist of two staff and three students*) and is "of questionable acceptability" if the collective takes determiner *one* or *another* (*One committee, appointed last year, has/?have not yet met*).

Both collective NPs and number-transparent quantificational NPs override the head's morphological singular at agreement. The override is possible at all because plural agreement requires only non-singular construal, not the precise atomic enumeration that singular *a(n)* and low cardinals demand. INDIVIDUATION — the speaker's construal of referents as discrete units, the variable that count grammar tracks — comes in different intensities. LOOSE count properties (agreement, *many/few*, plural-only forms) can be satisfied by individuation alone; TIGHT properties (singular *a(n)*, low cardinals) cannot, since they require selecting exactly one atom or enumerating a precise *n*-tuple. That asymmetry is why plural agreement can override a morphologically singular head while tight properties on the same NP follow ordinary count-noun behaviour. The override is selective in this way: it operates at the loose end of the cluster, not across it. Traditional grammar calls the agreement-side reflex NOTIONAL CONCORD. Reynolds (2026a) develops the loose/tight asymmetry as a precision-driven dissociation pattern across English count nouns and proposes a bidirectional link between morphosyntax and individuation as the underlying mechanism, where each cues the other and the override is

²A pilot COCA check (May 2026) across matched complements quantifies the contrast: plural-only forms show $\approx 0.2\%$ partitive uses; *plenty* $\approx 0.03\%$; *rest / remainder* show the reverse pattern at $\geq 97\%$ partitive (raw, higher after filtering idiomatic *for the rest of time*). Plural-only forms also reject *of [a N]* (0 hits across tested complements), indicating their *of*-complement is bare/non-determined rather than just non-partitive. Full counts and methodology in Appendix A.2; raw data in `paper/data/coca-pilot.md`.

the case where individuation outruns the head's morphological signal.

Where the non-singular construal is sourced differs in the two cases. In collectives, individuation is lexically encoded in the head: a singular collective denotes a unit whose members remain individuable, and the agreement relation accesses that member-level construal under predicates that license it. In number-transparent quantificational NPs, individuation is supplied by the *of*-complement: *a lot of tourists* sources non-singular construal from *tourists*; *a lot of money* sources low (non-count) construal from *money*; the construction projects the complement's count/non-count profile to agreement. The former is lexical-semantic accessibility of member individuation; the latter is constructional accessibility of the *of*-complement's individuation profile.

The two sources also differ in how stably the construal projects. The constructional source projects uniformly: in number-transparent QNs the override applies regardless of predicate type. The lexical-semantic source projects selectively: collective override requires the predicate to license member-level construal, which is why *The committee consists/*consist of two staff and three students* blocks the override even with the same singular collective head, and why singular-only determiners degrade plural override (*One committee, appointed last year, has/?have not yet met*).

Reynolds (2026a) distinguishes three levels of countability that don't always align: the ONTOLOGICAL (whether referents exist as bounded individuals in the world), the SEMANTIC (whether speakers construe referents as discrete units — individuation), and the MORPHOSYNTACTIC (whether nouns exhibit the grammatical count cluster: singular/plural contrast, *a(n)*, low cardinals, *many/few*, plural agreement, etc.). On that distinction, both override patterns sit at the interface of the second and third: the agreement relation accesses semantic individuation despite the head's morphosyntactic signal. The lexical-vs-constructional source axis is where individuation lives relative to the agreement relation in each case. Both routes are mediated accessibility under the umbrella schema; the difference is whether *P* (the individuation profile) is baked into the head's lexical semantics or imported from the *of*-complement.

CGEL itself notes that “the division between collective and number-transparent nouns is by no means sharply drawn” (Huddleston & Pullum, 2002, §18.2). Some nouns straddle. *Couple* has two senses: *couple*₁ (a man and woman in a relationship) is a regular collective; *couple*₂ (“approximately two”) is number-transparent. *Bunch* tilts collective for inanimate aggregates (*A bunch of flowers was presented to the teacher*) but transparent for groups of people (*A bunch of hooligans were seen leaving the premises*) (Huddleston & Pullum, 2002, §18.2).³

The schema accommodates both the paradigm and the transitional case, but they have different internal structure. Number-transparent NPs are class-like, with the override the default behaviour for class members. Collective NPs are gradient: the override probability depends on the head, the predicate, the determiner, the dialect register, and the construal the speaker has in mind. A theory of morphosyntactic transparency that ignored the difference between a class-based default override and a construal-driven optional one would lose what's interesting about both cases.

³A pilot COCA check (May 2026) supports the animate half (40 plural / 0 singular for *a bunch of people / kids*) but doesn't support the inanimate half: *a bunch of flowers was* and the CGEL example *A bunch of flowers was presented to the teacher* return 0 hits, and the two inanimate-bunch hits in subject position that did surface take plural rather than singular. The COCA evidence treats *bunch-as-subject* as transparent for animate aggregates and the inanimate-singular pattern as an analyst's example rather than a corpus regularity. Full counts: Appendix A.4.

2.3 COORDINATION AS A NON-INSTANCE

Coordination of singulars triggers plural agreement: *Bob and Jane are happy*. This looks like a transparency case at first (the conjuncts' number features percolate up to the coordinate construction and determine agreement), but it isn't. The schema requires that *X could have made P inaccessible*. The coordinate construction does no such hiding: there is no head whose morphological number competes with the conjuncts' features for agreement. The plural agreement is compositional addition ($1 + 1 = 2$), not mediation. Some coordinate NPs take singular agreement: *Bacon and eggs is my favourite breakfast*; *Wine and cheese is expected*; *Fish and chips is expensive these days*. These don't undermine the exclusion. In such cases agreement tracks a holistic construal of the coordinate phrase as a dish, package, event type, or conventional pairing. The singular feature is not a property of either conjunct that remains accessible through the coordinate construction; it is a construal of the coordinate phrase as a whole. Ordinary plural coordination and holistic singular coordination therefore differ in agreement semantics, but neither is transparent in the schema's sense. The former is compositional summation; the latter is whole-phrase unit construal.⁴ CGEL treats coordination-and-agreement under a separate heading (Huddleston & Pullum, 2002, §18.4), not alongside the override constructions of §18.2.

The clean distinction across §2's cases: in collective plural agreement, internal member plurality becomes accessible through a singular head; in number-transparent QNs, complement number becomes accessible through a quantificational head; in unit-construal coordination, the whole coordinate phrase is construed as singular and no internal property is being made transparent. The first two are mediated accessibility; the third is not.

The "*X could have made P inaccessible*" clause works as an independent diagnostic, not as an ad hoc filter. The test asks whether the construction-*type* admits a default-conceal variant: is there a version of the same construction-type, attested in the language, in which *X*'s internal structure makes *P* inaccessible to *R*? For English NPs, the head's morphological number determines agreement in the default case; the number-transparent class is the marked exception within the same NP construction-type. For NN compounds (§4), idiomatic and partly opaque compounds conceal compositional access; fully transparent compounds are at the other end of the same gradient. For transitive sentences (§3.2), frozen idioms like *kick the bucket* conceal parts' meanings; fully compositional transitives are at the trivial limit of the same gradient. For coordinate constructions, no such variant exists: the coordinate construction-type doesn't admit a version in which conjunct features fail to reach agreement, so coordination has no default-conceal backdrop. The plural agreement in *Bob and Jane are happy* is feature summation ($1 + 1 = 2$), not the marked side of a gradient that runs through concealment. The diagnostic therefore distinguishes mediator-types (those admitting a default-conceal variant, where the schema's accessibility question has bite) from compositional-summation types (where it doesn't). The schema is contrastive within mediator types: transparency is informative only where the same mediator type has opaque or partially opaque alternatives.

⁴Cases like *One and one is two* are an arithmetic-expression construction, not a meal/package coordination: *and* functions close to an addition operator, and the subject denotes an operation, equation, or expression. Not transparency either.

2.4 NUMBER-TRANSPARENT NPs AS A PROJECTIBILITY PROFILE

Section 1 set the diagnostic question: does identifying a candidate transparency profile let us predict behaviour? For number-transparent quantificational nouns the answer is yes, with the prediction running in two task settings.

For known CGEL members (*lot*, *plenty*, *lots*, *bags*, *heaps*, *loads*, *oodles*, *stacks*, *remainder*, *rest*, *number*, *couple*₂), lexical identity is sufficient input: knowing a noun is on this list licenses projection to the rest of the profile.

For candidate new members, class discovery uses two diagnostics independently observable from any of the projected behaviours: quantificational semantics (an indefinite quantity rather than a precise count or kind) and high-frequency occurrence as head of an *of*-complement construction. A candidate that satisfies both is then tested against the projected cluster:

- near-obligatory agreement override in the central cases, singular- or plural-direction depending on subgroup;⁵
- characteristic determiner patterns by subgroup (*lot* most often with *a*, also with *the* in totality uses; *plenty* typically with no determiner; the plural-only forms typically with no determiner; *remainder*, *rest* typically with *the*; *number*, *couple*₂ typically with *a*);
- partitive vs. non-partitive *of*-complement availability;
- behaviour under the fused-head construction;
- count/non-count percolation of the *of*-complement to the whole NP, with consequences for selection by predicates that are count- or non-count-sensitive;
- restricted premodification (*a whole/huge lot* but not *a remarkable plenty*).

The class's sub-group structure makes the cluster's predictions concrete. Table 1 displays the per-noun pattern for the central diagnostics; the entries are CGEL-coded rather than corpus-coded, and the within-domain audit a fuller study would supply is left as the within-domain check §2.4's falsification condition references.

Table 1. Number-transparent quantificational nouns: behavioural matrix (CGEL-based pilot).

Sub-group / Noun	Determiner	Partitive <i>of N</i>	Non-partitive <i>of N</i>	Override direction
Singular, <i>a</i> -determined				
<i>lot</i>	<i>a</i> ; also <i>the</i> (totality)	□ (<i>a lot of the money</i>)	□ (<i>a lot of money</i>)	sg → pl with plural complement
<i>number</i>	<i>a</i>	□ (plural oblique)	□ (plural oblique)	sg → pl with plural complement

⁵A pilot COCA check (May 2026) supports the near-categorical pattern: predicted-direction agreement runs 96–100% across nine tested QN+complement combinations after filtering. *A lot of money* required parse-shift filtering (10 of 13 raw plural-agreement hits were inside relative clauses modifying a plural head, e.g. *people who earn a lot of money are successful*; full filtering in Appendix A.3). Other small-count counter-direction cells haven't been spot-checked but the headline pattern doesn't depend on filtering them. Full counts and methodology: Appendix A.3.

Sub-group / Noun	Determiner	Partitive of <i>N</i>	Non-partitive of <i>N</i>	Override direction
<i>couple</i> ₂	<i>a</i>	□ (plural oblique)	□ (plural oblique)	sg → pl with plural complement
Singular, Ø-determined <i>plenty</i>	Ø	not typical	□ (<i>plenty of money</i>)	sg → pl with plural complement
Plural, Ø-determined <i>lots, bags, heaps, loads, oodles, stacks</i>	Ø	□ (COCA pilot: ≈ 0.2%, §2.1)	□	pl → sg with non-count complement
Singular, <i>the</i> -determined <i>remainder</i>	<i>the</i>	□ (<i>the remainder of the time</i>)	□ (<i>*the remainder of time</i>)	sg (default)
<i>rest</i>	<i>the</i>	□ (<i>the rest of the time</i>)	□ (<i>*the rest of time</i>)	sg (default)
Boundary <i>majority, minority</i>	<i>a</i>	□	?	variable
<i>bunch</i>	<i>a</i>	□	□	variable

Override applies in the central cases for all class members; boundary cases show variable behaviour (§2.2). Fused-head behaviour and count/non-count percolation are predicted to apply across class members but aren't per-noun-audited here. Premodification is restricted across the class but with sub-group-specific patterns (*a whole/huge lot*; minimal modification with *plenty*; fewer attested premodifiers with the plural-only forms).

Each sub-group covaries: knowing a candidate noun's determiner pattern and partitive availability narrows the sub-group, which then licenses the override-direction prediction (and vice versa). The boundary cases have weaker covariance, which is why §2.2's gradient framing applies to them.

Figure 1 visualises the override-direction data from the COCA pilot. Predicted-direction agreement runs 96–100% across nine tested QN+complement combinations after filtering. The override is essentially categorical; small-*n* cells (*plenty of money*, *the rest of the money*, *the rest of the people*) have wide confidence intervals even at 100% point estimates.

This two-task structure keeps the prediction genuinely projective for the new-member case: inputs are quantificational semantics and *of*-complement-head distribution; the cluster of agreement-override / determiner-frame / partitive / fused-head / count-or-non-count / premodification behaviours is then testable on the candidate. Boundary cases like *majority* and *minority* (§2.2) and the dual-sense *couple*₁ (collective) versus *couple*₂ (number-transparent) show graded membership at the edges;




Figure 1: Override-direction agreement for nine number-transparent QN + complement combinations in COCA. Each row is a QN+complement; the x-axis is the proportion of finite-verb agreement in the plural direction (0 = all singular, 1 = all plural). Predicted-singular cases (coral) cluster at the left; predicted-plural cases (blue) cluster at the right. Horizontal bars are 95% Wilson binomial confidence intervals; n is the total finite-verb count for that subject. *A lot of money* counter-direction count is filtered for parse-shifted hits (relative-clause modification of a plural-noun head); see footnote and paper/data/coca-pilot.md. Data: COCA pilot, May 2026.

the predictions hold in attenuated form, not as crisp class-conditional facts. What distinguishes prediction from redescription: the cluster supports inductive extension to *new* members (boundary cases like *majority/minority*; gradient cases like *bunch* in §2.2) that aren't built into the cluster's defining criteria.

The more demanding test is what the profile lets us predict beyond the in-grammar cluster. Three off-list projections are plausible candidates. *Diachronic*: members should enter the class through grammaticalisation or sense-split rather than through arbitrary lexical addition, with recent boundary additions (*majority*, *minority*) and the sense-split in *couple* matching that pattern. *Processing*: if the override is grammaticalised rather than improvised, sentences with number-transparent QNs shouldn't carry a comprehension cost relative to ordinary plural-subject sentences in adult speakers; no-cost or low-cost override is the prediction, not a parsing penalty for the surface mismatch. *Cross-linguistic*: the comparative concept of mediated accessibility (§6) licenses asking whether languages with analogous partitive-quantificational configurations exhibit comparable cluster structure (French *beaucoup de*, Spanish *un montón de*); the descriptive categories needn't be equated for the cluster shape to be compared. None of these is tested in the present paper. Their availability as predictions, rather than restatements of the input criteria, is what supports treating the class as a category in §7.1's stronger sense — a maintained-and-projectible cluster rather than just a distributional grouping.

The cluster of behaviours listed above makes number-transparent quantificational nouns the strongest candidate PROJECTIBILITY PROFILE examined here in the §1 sense: an observable base (lexical class membership for known members; quantificational semantics plus *of*-complement-head distribution for candidate members) supporting non-trivial projection to a structured set of targets, stabilised by conventional lexical inheritance. The profile is a candidate to pass Slater's (2015) broad stability-plus-projection diagnostic, with the within-domain check still empirically open. The falsification condition is concrete: if the cluster's predicted behaviours failed to track class membership in a substantial corpus or experimental study — if a noun satisfying the input diagnostics didn't show the predicted agreement, determiner, partitive, fused-head, count/non-count, and premodification pattern — the diagnostic would fail. The within-domain check is empirically open in this paper. *Majority* and *minority* (boundary instability) and *couple* (sense-split) reflect the conventional lexical inheritance that stabilises the category — qua linguistic class, the grouping is descriptively defined; qua category in §7.1's stronger sense, it is maintained-and-projectible. The count/non-count percolation behaviour the category licenses participates in the broader count cluster that Reynolds (2026a) analyses as a Khalidian causal cluster, with conventional individuation as a central causal node; on that analysis the percolation is embedded in causal architecture rather than free-floating convention, while Boyd-style homeostatic feedback for the count cluster as a whole is treated as a research hypothesis pending experimental perturbation evidence. Whether this profile earns the local-kind label as a downstream verdict is taken up in §7. The form-meaning accessibility family of §§3–5, with three within-domain profiles stabilised by different sources (constructional inheritance, usage entrenchment, processing architecture), is a separate research site where Khalidi-style refinement might apply if shared-item cross-modal evidence supports it.

2.5 TRANSPARENT FREE RELATIVES

A second morphosyntactic-transparency case fits the schema. Transparent free relatives (TFRs) are a sub-type of fused-relative construction (Huddleston & Pullum, 2002, ch. 12, §6.3) in which the matrix accesses properties of an embedded predicate (the NUCLEUS) through *what*: *what they call*

ecstatic, what we regard as essential, what seems to be a tourist. *What* is a fused-relative pronoun, morphologically and syntactically a singular indefinite; the matrix accesses features of the nucleus rather than of *what*. The construction has been examined extensively in the formal-syntactic literature, with three main competing analyses: a parenthetical-with-backward-deletion treatment (Schelthout et al., 2003; Wilder, 1999), a unified analysis on which standard and transparent free relatives share a configurational structure (Grosu, 2003, 2016), and a graft analysis in which a single constituent is shared between matrix and relative clauses (van Riemsdijk, 2006). The present treatment follows Kim's (2026, ch. 6.8) construction-grammar analysis and draws on the corpus evidence summarised in Reynolds (2026c).

Two *R*-relations dissociate empirically, and they have different verb-class profiles. AGREEMENT TRANSPARENCY matches matrix-verb agreement to the number profile of the nucleus and is licensed across the full inventory of TFR-licensing verbs, including the *seem/appear* class: *what appear to be trousers are on the chair* (plural agreement on plural nucleus *trousers* despite singular *what*; corpus-attested form, see (Reynolds, 2026c)). Attributional verbs license the same: *what they call essential workers are exhausted*. CATEGORY TRANSPARENCY is the narrower relation: the matrix's category requirement is mediated through the category of the nucleus, so a non-NP nucleus licenses the TFR in a non-NP slot. *He was what they call dimwitted* has an AP nucleus (*dimwitted*) and the TFR in predicative-AP position. This narrower relation is licensed only by the attributional sub-group of TFR-licensing verbs (*call, consider, describe as, regard as, take to be*); the *seem/appear* class shows agreement transparency only.⁶ Schema-wise, *X* is the TFR construction (or a sub-type of it); *P* is a property of the nucleus; *R* varies, and *R*-relativity is what makes the dissociation visible.

A working corpus check supports the asymmetry. AP nuclei are productive with attributional verbs (estimated hundreds of hits in COCA for *call* and *consider* alone, with sampled verification) but absent with *seem* and *appear* in a 178-hit exhaustive search; *take to be* patterns with the attributional class despite carrying an overt copula, indicating the constraint is on argument structure (the predicative complement's subject identified with the matrix object) rather than the presence of *be* (Reynolds, 2026c). Within the attributional class, PP nuclei are restricted to characterizing PPs (*what she'd call in love with Sam*), not locatives, suggesting a further sub-type stratification on the nucleus's predicative profile.

The projectibility analysis runs in two task settings, parallel to §2.4. For KNOWN TFR-licensing verbs, the verb's argument-structure type (subject identified with matrix subject vs. matrix object) is sufficient input: object-orientation is the diagnostic for the attributional daughter and predicts category transparency in addition to the agreement transparency licensed by the parent. For CANDIDATE verbs whose TFR-licensing status is uncertain (*prove to be, turn out to be, happen to be* are open cases), the same argument-structure diagnostic plus the AP-nucleus test isolates daughter membership without verb-by-verb stipulation. Two off-list projections are available. DIACHRONIC: the attributional class should show stable membership over time rather than free recruitment from the broader characterizing-verb inventory, with *take to be*'s occupying a small grammaticalisation pathway as the kind of marginal addition the cluster predicts. CROSS-LINGUISTIC: languages with parallel matrix-licensing constructions should exhibit comparable parent/daughter splits if the constraint

⁶A reader scanning CGEL for fused-relative examples with *seem* or *appear* will find NP-nucleus cases, since CGEL ch. 12 §6.3 treats fused relatives generally rather than the TFR sub-type; those examples are agreement-transparency cases under the present analysis, consistent with the verb-class restriction holding only for category transparency.

is on argument-structure configuration rather than on language-particular lexical content. Neither is tested in the present paper. Their availability as predictions, beyond the input criteria, supports treating the daughter as a candidate projectibility profile rather than as a distributional grouping.

The verb-class restriction itself admits two readings, both of which predict the asymmetry; distinguishing them is a research question the present paper doesn't settle. On a METALINGUISTIC reading, the attributional construction is quotative: *call* and *consider* attribute a label rather than directly predicating, and the attribution context relaxes the matrix's selectional pressure on the nucleus's category in a way that ordinary predication doesn't. On a MENTION-MEDIATED reading, the nucleus is treated as a use-mention hybrid; the construction marks the nucleus as a citation form whose category requirement is suspended. The two differ on what licenses the relaxation: speaker-attribution versus citation-marking. They make the same first-pass predictions; cancellability and projection tests on the resulting CIs would be diagnostic.

TFRs are a second morphosyntactic projectibility profile, with the cluster of agreement, category, and verb-class facts running off constructional sub-type membership rather than off the lexical-class membership that stabilises §2.1. The stabilising structure is constructional inheritance with a verb-class restriction at the daughter level. The constructional analysis is taken up in §3.5; full empirical and analytical development is in joint work with Jong-Bok Kim, in preparation.

2.6 KOREAN

[to be completed]

2.7 BRIDGE

The thread §2 traces is one shape of mediated accessibility: semantic individuation made available to morphosyntactic agreement, with category accessibility added by the TFR case. The lexical-vs-constructional source axis (§§2.2, 2.5) sorts cases within that thread; §2.2's three-level distinction (ontological discreteness, semantic individuation, morphosyntactic countability) places §2's territory at the interface of the second and third, with the source axis locating where individuation sits relative to the agreement relation in each case. Number-transparent quantificational nouns are the lexical-source paradigm, with collective NPs as transitional cases, coordination as a non-instance, transparent free relatives as the constructional-source paradigm, and Korean (per §2.6) extending the pattern to argument-role and information-structure relations. The next sections follow a different exit point off semantic individuation: §§3–5 take up form-meaning accessibility, where the relation *R* is interpretation, rated transparency, and processing rather than agreement. The umbrella schema covers both routes; the *R*-shift is what distinguishes them.

3 CONSTRUCTIONAL TRANSPARENCY

Section 2 treated agreement transparency: the mediator is a syntactic construction, the *of*-complement's number remains accessible despite the head's morphological number, and the relation is agreement. Section 3 turns to a different relation. Constructions also vary in how much of their internal form-meaning organisation remains accessible to speakers, learners, and analysts. This is the first of three sections (§§3–5) on the form-meaning accessibility family: §3 on the theoretical compositionality gradient, §4 on rated compound transparency, §5 on time-sensitive lexical processing. The three sections are within-domain profiles, distinct in items and methods; whether they cohere into a cross-domain projectibility cluster is the question §7 returns to, not a result the family is treated as

having delivered. §3.1 sets up the Construction Grammar framework. §3.2 treats compositionality as a gradient. §3.3 makes the section's central claim: constructional transparency is a network property, not a property of individual constructions in isolation. §3.4 connects §2's number-transparent NPs to the CxG framework. §3.5 returns to transparent free relatives (§2.5) under the inheritance lens. §3.6 looks ahead to §§4–5.

3.1 CONSTRUCTIONS AS FORM-MEANING PAIRINGS

Construction Grammar (henceforth CxG) treats constructions, learned form-meaning pairings, as the basic units of grammar (Goldberg, 2006). Constructions exist at all levels: morphemes (*pre-*, *-ing*); words (*avocado*, *anaconda*, *and*); complex words (*daredevil*, *shoo-in*); the regular-plural pattern [*N-s*]; filled idioms (*going great guns*, *give the Devil his due*); partly-filled idioms (*jog <someone's> memory*, *send <someone> to the cleaners*); Covariational Conditional (*The X-er the Y-er*, e.g., *the more you think about it, the less you understand*); Ditransitive (*Subj V Obj1 Obj2*); Passive (Goldberg, 2006, 5, Table 1.1). The list is heterogeneous; the unifying claim is that grammatical knowledge consists of these pairings, not of derivations applied to abstract underlying structures.

Argument structure isn't provided by verbs alone but by argument-structure constructions that combine with verbs. Goldberg's worked example uses *slice* across five argument-structure constructions: transitive (*He sliced the bread*), caused motion (*Pat sliced the carrots into the salad*), ditransitive (*Pat sliced Chris a piece of pie*), way construction (*Emeril sliced and diced his way to stardom*), resultative (*Pat sliced the box open*) (Goldberg, 2006, pp. 6–7). *Slice* contributes a constant lexical meaning across all five. The argument-structure semantics (someone acting on something; causing something to move; intending someone to receive something; moving despite obstacles; causing change of state) comes from the constructions.

Two further commitments follow. CxG is *USAGE-BASED*: patterns are stored as constructions if some aspect is non-predictable, or if they're fully predictable but sufficiently frequent (Goldberg, 2006, p. 5). The framework is non-derivational: “what you see is what you get,” with no underlying levels, no movement, and no phonologically null elements (Goldberg, 2006, p. 11). Both commitments matter for the transparency question, since they determine what counts as a construction's “parts” and what counts as the construction's contribution.

In the umbrella schema's vocabulary: *X* (the mediator) is the construction; *P* (what remains accessible) is the meanings of the parts and their compositional contribution to the whole; *R* (the relation) is interpretation, learning, and analogical generalisation. The relation is broader than §2's agreement, and the accessibility is gradient rather than class-like.

3.2 THE COMPOSITIONALITY GRADIENT

Compositionality is best treated not as a binary property but as a gradient. The working scale below discretises that gradient into four points for expository convenience. CxG itself presents compositionality as continuous; the four-point scheme is borrowed from the broader idiom-classification tradition (which distinguishes fully compositional from partly schematic and idiomatic constructions), not from CxG's own architecture, and is used here as a presentation aid rather than a CxG-internal claim.

Type	Transparency profile
Fully compositional pattern	parts and arrangement predict the interpretation
Partly schematic construction	construction contributes meaning, but parts remain interpretable
Idiomatically constrained construction	construction-specific meaning partly overrides compositional access
Frozen idiom	parts have little current interpretive accessibility

At the FULLY COMPOSITIONAL end, simple transitive sentences (*Pat sliced the bread*) have parts whose meanings predict the whole. The construction contributes argument linking, tense, and agreement, but the propositional content reduces to verb-meaning plus arguments. Fully compositional cases are the limiting endpoint of form-meaning transparency, not the paradigm. The transitive construction-type admits a default-conceal variant — frozen idioms like *kick the bucket* conceal compositional access — so any transitive sentence sits on a gradient where concealment is a real option for the construction-type; fully compositional sentences register as transparency at the trivial limit of that gradient. This is what differentiates them from §2.3’s coordination case: the coordinate construction-type admits no default-conceal variant, so coordination doesn’t sit on the gradient at all and the schema excludes it.

PARTLY SCHEMATIC CONSTRUCTIONS are where CxG does most of its empirical work. The Ditransitive (*Subj V Obj₁ Obj₂*) is the canonical example. *Liza bought Zach a book* and *Liza sent Stan a book* both inherit the construction’s meaning of intended transfer to an animate recipient, even though *buy* and *send* don’t lexically encode that meaning. The animacy constraint on the recipient explains the contrast between *Liza sent Stan a book* and ??*Liza sent storage a book* (Goldberg, 2006, p. 10). The construction contributes transfer-with-animate-recipient meaning; the verb contributes its lexical content. The two are interpretable independently and combine compositionally, but the combination is constructional, not just lexical.

IDIOMATICALLY CONSTRAINED CONSTRUCTIONS include partly-filled idioms (*jog <someone’s> memory*; *send <someone> to the cleaners*). The construction has a fixed lexical core whose meaning is metaphorical or non-literal, but the open slot is fillable by ordinary NPs and the surrounding syntax is regular. *Memory* in *jog memory* is interpretable as the same noun that appears elsewhere; *jog* contributes a metaphor of brief shaking. The construction-specific meaning (“cause to recall briefly”) partly overrides the compositional reading, but the parts remain visible. Schema-wise: the construction is the mediator; the parts’ meanings remain partly accessible; the relation is interpretation, but the interpretation is partly construction-specific.

FROZEN IDIOMS (*kick the bucket*) test the lower bound. The parts contribute nothing to the construction’s meaning of “die”; *kick* and *bucket* are dead lexically inside the construction even if they remain alive elsewhere. At this end of the gradient, *P* (the parts’ meanings) is no longer accessible through *X* (the construction) for *R* (interpretation). The construction passes the schema’s mediator-of-form requirement (it has internal structure that *could* have mediated access) but fails the accessibility-remains-through-it test. Frozen idioms mark where transparency in the umbrella sense bottoms out for the interpretation relation.

The gradient isn’t unique to argument-structure or idiomatic constructions; it runs through morphology too. Productive prefixation (*un-* in *unhappy*) sits near the fully compositional end. Partly

opaque derivatives (*business* from *busy*) sit near the idiomatic end: the historical compositional source is recoverable for analysts but not for ordinary speakers. The same gradient logic recurs in §4 (compound transparency) and §5 (priming-graded morphological transparency), though the domains and operationalisations differ.

3.3 CONSTRUCTIONAL TRANSPARENCY AS A NETWORK PROPERTY

§3.2 treated each point on the gradient as a property of an individual construction. The present paper develops a stronger claim that builds on Goldberg’s framework without being orthodoxy in it: constructional transparency also depends on the network in which a construction sits. Inheritance hierarchies link broad generalisations to subregularities to exceptional low-level constructions (Goldberg, 2006, p. 14). A construction’s parts and their compositional contribution are accessible partly because related constructions in the network make them so.

Goldberg’s worked example: the *What’s doing?* construction (*What’s that fly doing in my soup?!*) has a fixed form and connotes unexpectedness; it inherits from the Left Isolation construction, the Subject-Auxiliary Inversion construction, the Subject-Predicate construction, and the Verb-Phrase construction (Goldberg, 2006, p. 14). A learner who treats *What’s the X doing Y?!* as a construction in its own right does so partly because each of its inherited constructions has its own attestations and contributes a piece of the form-meaning organisation. Without the inherited constructions in the network, the learner would have less to work with.

This makes constructional transparency a network property in two distinct senses. **INTERPRETIVE TRANSPARENCY**: a construction’s compositional reading is accessible partly because related constructions provide compositional support. The Ditransitive’s transfer meaning is interpretable partly because the *for*-dative paraphrase (*Liza bought a book for Zach*) and the regular argument-structure constructions provide compositional context (Goldberg, 2006, p. 10). **LEARNABILITY TRANSPARENCY**: a construction is acquirable partly because its place in the inheritance network gives the learner multiple routes of access. Both senses are *R*-relativised in the umbrella schema: *P* remains accessible for the purposes of *R*, where *R* is interpretation in the first sense and learning or analogical generalisation in the second.

The network claim is offered as a hypothesis, not as established. If the hypothesis holds, constructions in dense regions of the constructicon would be more transparent for both *R*-values than constructions in sparse regions, and constructions whose family-resemblance relatives have low type frequency would be harder to interpret and learn. Goldberg (2006, p. 15) cites usage-based acquisition work (Diessel and Tomasello, Tomasello, Saffran and colleagues) as part of the broader case for usage-based acquisition; whether any of those studies bear specifically on transparency-as-network-property is itself an empirical question the present paper doesn’t settle. The form-meaning accessibility family of §§3–5 will return to the hypothesis: the corpus-derivable predictors that Bell and Schäfer use in §4 and the priming-magnitude modulations that Heyer and Kornishova report in §5 will be treated as compatible quantitative probes for the network-relative transparency that §3 frames theoretically, not as direct measurements of it.

The network claim also clarifies what frozen idioms mark. *Kick the bucket* has nearly no parts-to-whole compositional accessibility, but it’s recognisable as an idiom rather than as random word combinations. That recognition is partly supported by the existence of partly-idiomatic constructions one rung up the gradient: network position helps explain why the string is treated as a unit. The construction can also be stored holistically and accessed without compositional decomposition;

nothing in the network claim rules that out. Even at the bottom of the gradient, network position does some work, but the work is recognition-relative, not interpretation-relative.

3.4 NUMBER-TRANSPARENT NPs REVISITED

§2 treated CGEL number-transparent NPs as agreement-accessibility cases. From the CxG side, those same NPs can be analysed as constructions: form-meaning pairings whose form is the *number-transparent quantificational noun + of-complement* template and whose meaning includes the override specification on agreement. The two analyses trade off different formal properties. CGEL preserves a head-and-complement structure (with *number*, *lot*, etc. as heads) and treats the override as a property of the construction rather than as a separately stipulated mechanism — the term *percolation* appears only informally in a footnote (Huddleston & Pullum, 2002, §18.2 fn. 12). A CxG analysis bundles the form and the agreement specification into a single constructional schema; whether the schema’s internal representation preserves head-and-complement structure depends on the specific CxG variant. Sign-Based Construction Grammar preserves it explicitly; some Goldbergian variants leave the internal representation more abstract. The choice among CGEL, SBCG, and Goldbergian CxG is methodological. The umbrella-schema-level claim that X (the construction) mediates between P (the *of*-complement’s number) and R (agreement) is supported under any of these formal apparatuses. The present paper doesn’t develop the CxG-side analysis substantively; §3.4’s purpose is to flag the methodological choice and show that the umbrella schema sits above either formalism, not to commit to one.

This connects §2 and §3 in a useful way. Number-transparent NPs are constructional-transparency cases, but the relation R is agreement (a §2-style morphosyntactic relation), not interpretation (the §3-style relation). The umbrella schema is one frame; the senses partition by R . A construction can be transparent for one R and opaque for another: a frozen idiom like *kick the bucket* is opaque for interpretation but morphosyntactically regular (the verb agrees with its subject in the ordinary way). R -relativity is built into the umbrella; it’s a feature, not a bug. The schema asks whether P remains accessible *for the purposes of* R , and accessibility isn’t transitive across R -values.

The methodological consequence: choosing CGEL’s apparatus or CxG’s apparatus for any given case is a question about which formal vocabulary makes the empirical generalisations clean. It isn’t a question about which framework is “right” in some deeper sense. The umbrella schema sits above both: it states the relational claim that a transparency analysis has to satisfy, without committing to how the satisfaction is formally encoded.

A coherence point across the frameworks the paper engages: CGEL, SBCG, and Goldbergian CxG are all non-derivational, and none posits empty mediating elements (traces, copies) between a surface position and a hypothesised underlying one. Contemporary dependency-grammar accounts (da Cunha & Gibson, 2025; Gibson, 2025) take the same line on long-distance dependencies, arguing that direct head-to-dependent connections fit acceptability and processing data without empty mediators (cf. the §2 extraction-transparency footnote). The umbrella schema doesn’t adjudicate the dispute, but the frameworks engaged here all treat the mediator X as observable structure rather than as posited invisible mediation between observable layers. Where the schema asks whether P remains accessible through X for R , the candidate X is always something the analyst can point at.

3.5 TRANSPARENT FREE RELATIVES REVISITED

The TFR case introduced in §2.5 sits naturally in §3's frame. The agreement/category dissociation across the verb inventory is constructional transparency in a tight sense: a constructional sub-type (the attributional daughter) has access conditions that the parent construction doesn't, and the access condition tracks an inheritance relation rather than a per-item lexical property. A TFR-parent construction licenses agreement transparency across the full verb inventory; a TFR-attributional daughter, restricted to attributional verbs and inheriting from the parent, licenses category transparency in addition. Kim's (2026, ch. 6.8) construction-grammar treatment is the architectural anchor; Reynolds' (2026c) corpus check supplies the empirical content for the parent/daughter split. As §2.5 notes, the daughter's category-transparency licence has two competing readings: a metalinguistic-attribution account (the construction is quotative; *call* and *consider* attribute a label rather than directly predicating) and a mention-mediated account (the construction marks the nucleus as a citation form whose category requirement is suspended). Both predict the verb-class restriction; both are constructional in §3's sense, since both locate the licence at the daughter rather than in any individual lexical entry. The work §3.5 needs to do here is independent of which reading wins: the daughter's licence is a property of the construction's network position, not of the verbs that fill it.

The TFR case gives §3.3's network hypothesis a tractable instantiation, though not a confirmation: the daughter's predictions (verb-class restriction, PP-restriction to characterizing PPs per §2.5, the *take to be* placement that defeats a copula-as-blocker analysis) are stabilised by the construction's position in an inheritance network. They aren't attached to any particular verb's lexical entry, and they don't reduce to compositional consequences of the matrix's selectional requirements. *R*-relativity is also visible in the case: the same mediator *X* is transparent for one *R* (agreement, both daughters) and only partially transparent for another (category, daughter-only). Two *R*-relations under one constructional sub-type partition cleanly along an inheritance line. The CGEL-vs-CxG methodological choice §3.4 flagged for number-transparent NPs reappears here: a sign-based variant with explicit inheritance hierarchies (SBCG, the framework Kim deploys) makes the parent/daughter relation the central analytical move, while a Goldbergian variant absorbs the same relation into a network of family-resemblance constructions. The umbrella-schema-level claim survives either choice; the differential constructional licence sits above the formalism.

3.6 LOOKING AHEAD TO §§4–5

§3 has set up the theoretical frame for the form-meaning accessibility family: constructions as form-meaning pairings, compositionality as gradient, transparency as network property. Section 4 will give the quantitative measurement. Bell and Schäfer's (2016) models of compound transparency operationalise expectedness via corpus statistics and human ratings, providing empirical content for the partly-schematic and idiomatic regions of the gradient. Section 5 will give the processing measurement. Heyer and Kornishova (2018) show that semantic transparency modulates masked priming, but only at long stimulus-onset asynchronies, a finding that bears on the time-course of network access.

The three sections together motivate the form-meaning accessibility family, with three evidential windows that bear on it; whether they converge into a cross-domain projectibility cluster is what §7 will weigh, and §7 doesn't claim convergence delivered. Whether each within-domain profile passes Slater's (2015) SPC test, and whether Khalidi's (2018) causal-network refinement applies to any cross-domain reading, are questions §7 returns to. For now, §3 has done its job by setting the frame: con-

structional transparency is gradient and network-relative, and the next two sections give the family its quantitative and processing-experimental bone structure.

4 SEMANTIC TRANSPARENCY

Section 3 set up form-meaning transparency as a theoretical gradient and as a network property; the relation R was interpretation, learning, or generalisation. Section 4 narrows R to interpretation and turns to a quantitative case: rated transparency of English noun-noun compounds. This is the second of three windows on the form-meaning accessibility family, between §3's theoretical frame and §5's processing evidence. §4.1 lays out Bell and Schäfer's model. §4.2 connects it to the umbrella schema. §4.3 places referential transparency as the schema's boundary case, where X is a context rather than a form. §4.4 returns to the cluster question. §4.5 bridges to §5.

4.1 COMPOUND TRANSPARENCY: THE BELL AND SCHÄFER MODEL

Bell and Schäfer (2016) model the perceived semantic transparency of English noun-noun compounds as a function of *expectedness* in their internal semantic structure. A compound is rated as more transparent when its semantic structure is more expected, both in the predictability of its constituents in their positions and in the predictability of the semantic relation between the constituents given those constituents. Transparency is treated as a scalar variable rather than a binary or four-fold property. The empirical anchor is the database of Reddy and colleagues (2011): 30 ratings on a 0–5 scale for each of 90 two-part NN compounds, totalling 8,100 ratings (Appendix B.5).

Bell and Schäfer's models identify three predictors of higher rated transparency. First, *N₁ frequency*: compounds whose first noun occurs more frequently in the language overall are rated as more transparent. Second, *semantic-relation expectedness given N₁*: compounds whose semantic relation is more frequently associated with that particular N₁ are rated as more transparent. Third, *N₂ productivity*: compounds whose second noun occurs more often as the head of NN compounds are rated as more transparent. The first and third predictors are corpus-statistical; the second requires hand-coding of semantic relations using Levi's (1978) nine-predicate inventory (CAUSE, HAVE, MAKE, USE, BE, IN, FOR, FROM, ABOUT; Appendix B.1) and aggregation across compounds sharing each N₁.

The two headline conclusions: N₁ and N₂ contribute asymmetrically to whole-compound transparency, and all significant predictors of rated transparency also independently predict processing speed. The latter motivates Bell and Schäfer's tentative hypothesis that “perceived transparency may itself be a reflex of ease of processing” (Bell & Schäfer, 2016, p. 157). The hypothesis is offered as a direction for further work, not a result. Stress-pattern speculation about the IN/FOR positive associations is in Appendix B.3; the N₁/N₂ asymmetry detail is in Appendix B.4.

4.2 COMPOUND TRANSPARENCY IN THE UMBRELLA SCHEMA

The umbrella schema's variables map directly. X (the mediator) is the compound, e.g. *car-wash*, *jail-bird*, *hogwash*. P (what remains accessible) is the constituent meanings and their semantic relation. R (the relation) is interpretation: rated transparency operationalises how successfully the constituents' meanings reach interpretation through the compound. The schema's “remains accessible” requirement is gradient here, just as in §3: a compound's parts are accessible to varying degrees, with rated transparency as the empirical proxy.

Bell and Schäfer's three predictors fit naturally inside the schema. N₁ frequency, N₂ productivity, and semantic-relation expectedness given N₁ are properties of the network in which a compound sits, not properties of the compound in isolation. A compound is rated more transparent when its constituents and the relation between them are more conventional in the language. This is the network-relative view §3.3 framed theoretically, expressed quantitatively. Compound transparency isn't a property of the compound alone but of the compound's position in the construction-and-corpus.

One caveat about reading the network claim into Bell and Schäfer's results. Their predictors are corpus-statistical; the construction network is a theoretical posit. The two are compatible (frequency-driven entrenchment is the standard usage-based story for how the construction is shaped), but Bell and Schäfer don't directly measure construction-network properties. Their work is consistent with the network claim without testing it. Section 3's theoretical framing and §4's quantitative results are mutually compatible; that's the operative claim, not a stronger one.

4.3 REFERENTIAL TRANSPARENCY AS BOUNDARY CASE

The schema's referential boundary is independent of the form-meaning accessibility family. *Lois Lane believes Superman is brave* is a textbook intensional context; substituting *Clark Kent* for *Superman* yields a sentence whose truth-conditions may not be preserved (the *de re* reading allows substitution; the *de dicto* reading doesn't). This is referential transparency in the philosophical-semantics tradition, and the schema technically applies: *X* is the intensional context (the embedding under *believe*), *P* is the referent of the embedded singular term, *R* is substitution under co-reference.

What makes this the boundary case is that *X* is a *context* rather than a *linguistic form*. In all the cases §§2–3 worked through, *X* was a form whose internal structure could have mediated access to *P*: a quantificational-noun-of-complement construction; a compound; a partly-schematic argument-structure construction. An intensional context isn't a form in that sense. It's a configuration that creates substitution-opacity, and the *de re/de dicto* distinction is about whether substitution is allowed, not about whether form-internal compositional access is preserved.

The schema therefore stretches thinnest at this boundary. *P* is a referent, not a constituent meaning or a number feature; *R* is substitution-*salva-veritate*, not agreement or interpretation. The umbrella applies on a sufficiently abstract reading: referential transparency is a kind of mediated accessibility, where the mediator is the intensional embedding and the referent is what does or doesn't remain accessible. But the difference in the type of *X* is real and matters for the metaphysical conclusion. Referential transparency isn't part of the form-meaning accessibility family of §§3–5.

The treatment here is intentionally controlled. Referential transparency has a substantial philosophical literature whose canonical mid-century anchor is Quine (1956); the broader tradition (Frege, Kripke, Kaplan) would dominate the section if developed. The methodological point is that not every transparency relation in linguistics fits the form-meaning accessibility family, and the schema's *R*-relativity is what allows the paper to handle this case without forcing it in. The boundary marks a real heterogeneity in the umbrella, not a defect.

4.4 COMPOUND TRANSPARENCY AS A WITHIN-DOMAIN PROJECTIBILITY PROFILE

Compound transparency joins constructional compositionality (§3) as a within-domain projectibility profile on form-meaning accessibility. Class-membership talk doesn't apply at this level: there is no small, conventionally stable class of "transparent compounds" in the way there is for number-

transparent quantificational nouns. The story here is gradient: compounds are more or less transparent on a continuous scale, and the predictors of higher transparency are corpus-statistical and constructional.

Compound transparency licenses non-trivial within-domain projection. Knowing a compound's N₁ frequency, semantic-relation expectedness, and N₂ productivity lets us predict its rated transparency, with non-trivial fit in Bell and Schäfer's models. The stabilising source is usage-based entrenchment: the network of compounds and their semantic relations is shaped by frequency, productivity, and the kind of conventional patterns that emerge from repeated use. Falsification condition: if N₁ frequency, semantic-relation expectedness, and N₂ productivity failed to predict rated transparency in an extension or replication of Bell and Schäfer's models, the within-domain profile would fail. Whether the within-domain projection joins with §3's theoretical profile and §5's processing window into a cross-domain projectibility cluster on shared items is left to §7, and §7 doesn't take that result as delivered.

What §4 contributes that §3 didn't is empirical content. The compositionality gradient theorised in §3 is a measurable property whose predictors are corpus-derivable and whose values correlate with rated transparency. That correlation is what makes compound transparency a quantitative window onto the gradient. Whether the same predictors would also track processing latencies, which would strengthen the case for Khalidi-style causal-hierarchical refinement (without by itself establishing it), is a question for §5.

4.5 BRIDGE TO §5

§5 takes the third window on the family: time-sensitive lexical processing. Where Bell and Schäfer measure rated transparency, Heyer and Kornishova measure masked priming. Where rated transparency is a steady-state property, priming is a time-course property. Whether the windows of §3, §4, and §5 jointly support the kind of causal-hierarchical projection that Khalidi's framework formalises is the question §7 will return to.

5 PROCESSING TRANSPARENCY

Section 4 measured rated transparency at the level of whole compounds; the rating is a steady-state property of the compound, given by speakers' explicit judgements. Section 5 turns to a time-course property: how does semantic transparency modulate the processing of morphologically complex words during recognition? Heyer and Kornishova's (2018) masked-priming study is the empirical anchor. Their finding, summarised in their title's three-dot ellipsis: semantic transparency affects morphological priming, but only at long stimulus-onset asynchronies. §5.1 lays out the design and findings. §5.2 unpacks the serial-processing interpretation. §5.3 places it against a rival camp. §5.4 returns to the cluster question. §5.5 closes.

5.1 THE HEYER AND KORNISHOVA STUDY

Heyer and Kornishova (2018) ran two parallel masked-priming experiments, one with English *-ness* nominalisations and one with Russian *-ost'* nominalisations. The two suffixes are deadjectival and produce abstract nouns denoting "state of being *X*" (e.g., *dark* + *ness* = *darkness*; *bledn-* + *-ost'* = *blednost'* "paleness"). Targets in both experiments were the adjective bases (e.g., *pale*, *dark*); primes were the corresponding nominalisation, a simplex unrelated control, or the adjective as identity prime.

Each experiment crossed two stimulus-onset asynchronies (SOAs), short and long, with transparency as a continuous variable rated by separate groups of speakers (Appendix C).

Transparency was rated on a 1–7 scale: *business* (rated 3.17) at the opaque end is no longer interpretable as “state of being busy”; *falsehood* (rated 5.75) is fully derivable from *false*; items spanned the range continuously.

The headline result is SOA-conditioned. At short SOA in both languages, priming is constant across the transparency scale: morphologically complex primes are decomposed into stems regardless of whether the stem still contributes to the meaning of the whole. At long SOA, priming becomes transparency-modulated: more transparent prime-target pairs (e.g., *palehood-pale*) yield larger priming effects than less transparent pairs (e.g., *business-busy*). The three-way interaction (Prime Type × Transparency × SOA) is significant in English and marginally so in Russian, with the same direction in both languages (Heyer & Kornishova, 2018, pp. 1117–1119). Statistical detail is in Appendix C.3.

5.2 THE SERIAL-PROCESSING INTERPRETATION

Heyer and Kornishova read the data as evidence for a serial decomposition process. Morphologically complex words are first decomposed into stem-plus-affix on the basis of orthographic structure (*morpho-orthographic decomposition*), with semantic information entering only at a later stage. The first stage is semantically blind: it parses *palehood* into *pale* + *ness* regardless of whether the parse yields a coherent semantic interpretation, and would parse *corner* into *corn* + *er* even though the parse fails semantically (Heyer & Kornishova, 2018, p. 1112).

This is consistent with Bell and Schäfer’s (2016) §4 hypothesis that “perceived transparency may itself be a reflex of ease of processing.” Within the narrow Heyer-and-Kornishova-style derived-word masked-priming design, the pattern is that fast and shallow processing windows don’t show transparency modulation, while longer processing windows do. The “fast-and-shallow versus slow-and-deep” framing applies to that paradigm specifically; whether it extends to other tasks (whole-word reading, lexical decision without priming, comprehension) is a separate empirical question. The two studies are compatible with a shared form-meaning accessibility hypothesis, though they involve different items (NN compounds versus derivational nominalisations) and different methods (ratings/modelling versus masked priming) and so don’t test the same processing chain.

5.3 THE RIVAL PARALLEL-PROCESSING POSITION

The serial reading is contested. Feldman and colleagues’ parallel-immediate position argues that morpho-semantic information influences recognition from the earliest stages. Feldman et al.’s (2015) time-course analysis presses the diagnostic question (*must analysis of meaning follow analysis of form?*) with a same-target design — each target is compared after semantically similar and semantically dissimilar primes within subjects — and SOA modelled as continuous from 34 to 100 ms rather than as binned categorical levels. They report that semantic similarity affects recognition latencies even at SOAs of 34 ms and 48 ms, shorter than Heyer and Kornishova’s *short* windows of 39 ms (English) and 33 ms (Russian). On the parallel-immediate reading, the form-then-meaning ordering doesn’t survive that finding. The two designs aren’t head-to-head: Heyer and Kornishova’s short-SOA result is the absence of transparency modulation across an item set rather than the absence of semantic effects on a same-target comparison, and the disagreement is whether semantic information is constitutively present at masked-priming windows or emerges only after morpho-orthographic decomposition. The literature is mixed enough (the meta-table on p. 1121 of Heyer and Kornishova

surveys 35 prior masked-priming studies) that the present paper takes the serial reading as Heyer and Kornishova's reading, not as a settled disciplinary verdict. The cross-modal probe in Appendix D would generate within-items evidence relevant to the disagreement.

5.4 WHAT §5 CONTRIBUTES TO THE CLUSTER

§5 contributes the time-course dimension to the form-meaning accessibility family. §3 framed compositional transparency as a gradient; §4 made it quantitative through corpus predictors and ratings; §5 makes it temporally articulated. The family's evidential support is now: a theoretical scale, a steady-state quantitative measure, and a time-course measure. The three windows are operationally distinct but conceptually related: they're different probes for the same underlying form-meaning accessibility, taken at different stages of speakers' interaction with the relevant items. Falsification condition for the within-domain priming profile: if rated semantic transparency failed to modulate masked-priming magnitude at long SOAs in an extension or replication of Heyer and Kornishova's design, the profile would fail.

The relationship between §4's ratings and §5's priming deserves a careful word. Bell and Schäfer's compounds and Heyer and Kornishova's nominalisations are different items in different morphological domains; their results don't compose into a single transitive empirical chain (corpus predictors → ratings → priming on the same items). What they jointly support is the cluster hypothesis: rated transparency and priming-graded transparency both behave as the form-meaning accessibility hypothesis predicts, in a way that distinguishes them from idiosyncratic post-hoc descriptive labels. They're compatible probes for the same construct, not direct measurements of one another. Whether the cluster hypothesis survives the within-items test — corpus predictors, ratings, priming magnitude, and interpretation accuracy aligning on shared items as joint reflexes of underlying representational structure — is the empirical question §D sketches as a probe design. Until that probe is run, the family is three within-domain profiles plus a research hypothesis, not a delivered cross-domain cluster.

5.5 CLUSTER STATUS AND BRIDGE TO §6

§§3–5 give three within-domain projectibility profiles, each stabilised by a different source: a construction-level profile (§3, stabilised by constructional inheritance) predicts compositional interpretability via inheritance and network position; a compound-level profile (§4, stabilised by usage-based entrenchment) predicts rated transparency from corpus statistics; a derivationally-complex-word-level profile (§5, stabilised by processing architecture) predicts long-SOA priming magnitude from rated transparency. Whether they jointly license a cross-domain projectibility cluster on shared items, as opposed to remaining three within-domain profiles loosely related by analytical theme, is the question §7 takes up. Acquisition trajectory is a plausible further extension, consistent with the network claim of §3.3, but the present sections don't directly ground it. Khalidi's (2018) causal-network refinement applies at the cross-domain level only if shared-item evidence supports a causal-hierarchical reading, which the present paper doesn't supply. §5's contribution is to add a time-course profile to the within-domain evidence, not to establish cross-domain projection by itself.

§6 turns from form-meaning accessibility back to the comparative question raised in §1: how does the umbrella schema constrain crosslinguistic generalisation about transparency? Haspelmath's (2010) descriptive-versus-comparative-concept distinction is the right discipline for that question, and §6 applies it to the morphosyntactic and form-meaning cases the paper has worked through.

6 TYPOLOGICAL TRANSPARENCY

Section 6 is methodological, not substantive. The previous sections worked through transparency cases in English (§2 number-transparent NPs; §3 CxG examples; §4 compound transparency; §5 *-ness* nominalisations) and in Russian (§5 *-ost'* nominalisations), with a Korean stub (§2.6) reserved for co-authored work. Section 6 asks how the paper should handle the move from these language-particular cases to any cross-linguistic claim about transparency. Haspelmath's (2010) descriptive-versus-comparative-concept distinction is the right discipline for the question, and §6 applies it to the umbrella schema.

6.1 HASPELMATH'S DISTINCTION

Haspelmath (2010) argues that crosslinguistic comparison shouldn't assume universally available categories like ADJECTIVE, PASSIVE, ACCUSATIVE, or SUBJECT. Each language has its own descriptive categories, defined by the language-particular distributional and morphosyntactic criteria that make the categories work in that language (Haspelmath, 2010, p. 664). The criteria differ from language to language: what counts as a passive in one language isn't necessarily what counts as a passive in another, so descriptive categories can't be equated across languages.

Comparative concepts, by contrast, are analyst-created tools designed for the purpose of crosslinguistic comparison. They aren't part of any language system and aren't psychologically real; they can't be right or wrong, only more or less well suited to the comparative task (Haspelmath, 2010, p. 665). Comparative concepts have to be universally applicable, which means they have to be built from universally applicable building blocks: conceptual-semantic concepts (such as notions of agency or possession), general formal concepts (such as "phonological word"), and other comparative concepts.

Haspelmath calls the alternative position categorial universalism: the assumption that there is a substantial set of crosslinguistic categories from which languages select (Haspelmath, 2010, p. 663). Newmeyer's (2007) *Linguistic typology requires crosslinguistic formal categories* is the canonical statement of the position Haspelmath rejects: typological generalisations, on Newmeyer's argument, can't be made through purely conceptual-semantic comparative concepts because such generalisations must reference the specific form in which the concepts are realised (Newmeyer, 2007, p. 136). Haspelmath (2010, p. 666) grants the realisation point but argues that the relevant formal references are themselves comparative concepts (e.g. "adposition", "zero coding") definable in conceptual-semantic and general-formal terms, not crosslinguistic descriptive categories. Categorical particularism, his preferred position, holds that descriptive categories are language-particular and that comparison happens through analyst-defined comparative concepts that stand in many-to-many relations with descriptive categories. The position has been developed in typological work by Dryer, Croft, Lazard, and Cristofaro.

The methodological consequence is straightforward. A typological generalisation isn't a claim about a category instantiated in every language. It's a claim about how analyst-defined comparative concepts are realised in language-particular ways. The generalisation that nominative case-marking is widespread across languages doesn't entail that "nominative" picks out the same descriptive category in any two languages.

6.2 THE UMBRELLA SCHEMA AS A COMPARATIVE CONCEPT

The umbrella schema introduced in §1 is a comparative concept. It's built from conceptual-semantic notions (mediation, accessibility, property, relation) that are universally applicable, plus relational variables (X , P , R) whose values are filled in language-particularly. The schema's question (*X is transparent w.r.t. P for R when P remains accessible through X for the purposes of R*) is something the analyst asks of any language; it isn't a claim that any particular language has a descriptive category corresponding to TRANSPARENCY.

The building blocks deserve scrutiny against Haspelmath's standard. Haspelmath (2010, p. 665) restricts comparative concepts to conceptual-semantic notions, general formal notions, and other comparative concepts. The schema's primitives all fit. *Property* (P) and *relation* (R) are conceptual-semantic primitives common to descriptive grammar; they don't presuppose any language-particular grammatical content. *Form* (X) is a general formal notion: any constituent or configuration the analyst can identify in a language counts as a candidate X . *Accessible through* is the contested slot. *Accessibility* has discipline-specific senses (discourse-pragmatic salience, psycholinguistic retrievability, formal-semantic substitution-licensing); they're kin but not identical, and the paper doesn't claim they reduce to one another. The schema's intended sense is more abstract: P is accessible to R through X when R tracks P across X rather than tracking some property of X itself. That definition is conceptual-semantic — a relational predicate between a property and a relation, mediated by a form — without psycholinguistic commitments. The case studies operationalise this abstract sense in different ways (agreement targeting in §2, constructional interpretability in §3, rated transparency in §4, priming magnitude in §5), but the comparative concept lives at the abstract level. *Mediation* is the schema's other central notion: a form whose internal structure could have hidden the property and doesn't, tested by the contrast between the paradigm cases (§§2.1, 2.2, 2.5) and the non-instance (§2.3 coordination). Both building blocks are general enough that any language admits the schema's question, which is what Haspelmath's standard requires.

Each value of the schema's variables is language-particular. In English, X is the *number-transparent quantificational noun + of-complement* construction (a CGEL descriptive category); P is the number/count profile of the complement; R is subject-verb agreement. In Korean (per §2.6, pending in the co-authored extension), X may be a layered case-marking pattern or a scrambled word order; P may be argument-role assignment or information-structural status; R may be argument-role recovery or scope. The variable bindings are different. The comparative concept MEDIATED ACCESSIBILITY is the same.

This is what Haspelmath's framework requires: cross-linguistic claims are made through comparative concepts whose realisations differ from language to language. The English and Korean cases aren't "instances of the same category." They're different descriptive categories that can be compared through the comparative concept of mediated accessibility, with different X , P , and R values.

The form-meaning accessibility family of §§3–5 isn't itself a comparative concept. The schema (mediated accessibility, with its $X/P/R$ variables) is the comparative concept; the family is a set of language-particular phenomena and theoretical/experimental measurements compared *through* the schema. CxG constructional transparency, compound rated transparency, and morphological priming are different operational windows on form-meaning accessibility (§§3–5 worked them through in English with one Russian extension): a theoretical profile, a measurement, and an experimental effect. They're intended to bear on the same comparative-concept gradient. The cross-linguistic move

at the comparative-concept level doesn't commit the paper to crosslinguistic descriptive categories.

6.3 WORKED EXAMPLE: DESCRIPTIVE CATEGORY AND COMPARATIVE CONCEPT

Consider the typological question: do all languages have number-transparent constructions? The categorial-universalist answer is to look for English-style number-transparent NPs in every language. The categorial-particularist answer is different. NUMBER-TRANSPARENT QUANTIFICATIONAL NOUNS is a CGEL English descriptive category, defined by the override behaviour (near-obligatory in the central cases), the small and relatively stable membership, the characteristic determiner patterns, and the partitive/non-partitive *of*-complement restrictions worked through in §2.1. That category isn't what crosslinguistic comparison tracks; it's what the comparative concept *organises*.

What the comparative concept of mediated accessibility lets us ask is different. Across languages, which configurations let an internal property remain accessible to a grammatical relation that the surface form might otherwise have hidden? English number-transparent NPs answer the question for the relation R = subject-verb agreement, with P = number/count profile of the *of*-complement. The comparative concept doesn't predict that any other language has matching descriptive categories. It predicts that the analyst's question is askable of any language, and that the answers will differ.

The cross-linguistic extensions the paper has on hand are limited. §5's English *-ness* / Russian *-ost'* parallel is one: two languages, two descriptive categories (deadjectival nominalisation in English; deadjectival nominalisation in Russian), but the same comparative-concept question asked of both, with parallel answers. The Korean case that §2.6 stubs is the most consequential extension, deferred for separate co-authored development. Without those Korean answers, this section's worked example carries the methodological move with English alone; the cross-linguistic claim is a structural one (the comparative concept lives above the descriptive category, whatever the language) rather than a fully demonstrated comparative one.

This is the right level of crosslinguistic claim. It doesn't require positing a universal category TRANSPARENCY or NUMBER-TRANSPARENT. It doesn't require claiming that any two languages have "instances of the same construction." It claims, more modestly, that the comparative concept of mediated accessibility can be applied to any language and that the answers it elicits are language-particular categorial content. That modesty is what categorial particularism asks of comparative work.

6.4 THE REVIEWER OBJECTION AND A REPLY

A reviewer might object that the umbrella schema looks like a universal category smuggled in as a comparative concept. The paper holds that TRANSPARENCY is a family of senses rather than a single kind, but the umbrella schema applies across English, Korean, and the form-meaning family of §§3–5; is that not exactly what categorial universalism claims?

The reply is that the schema is a comparative concept by design. It's an analyst's question (does P remain accessible through X for R in this language?), not a claim that any language has a descriptive category called TRANSPARENCY. Each language answers the question differently or not at all. The schema's universal applicability is the universal applicability of the *question*, not of any *answer*. Haspelmath's (2010, p. 665) criterion for a good comparative concept (it can be more or less well suited to the task of permitting crosslinguistic comparison) is exactly the standard the umbrella schema is meant to meet.

Newmeyer's (2007) specific argument deserves a direct response. Newmeyer holds that typological generalisations have to reference language-particular formal realisations, not just conceptual-semantic concepts, and that this commits typology to crosslinguistic formal categories. The umbrella schema's variables address this without conceding the categorial-universalist conclusion. The schema's *X* slot is filled language-particularly: the English *X* is the *number-transparent quantificational noun + of-complement* construction (a CGEL English descriptive category); the English TFR *X* is the fused-relative construction with attributional-verb licensing (a CGEL/CxG English category); the Korean *X* (per §2.6) is whatever Korean construction the analyst identifies. The schema's universal applicability is at the analyst-question level; the formal realisations Newmeyer demands enter at the variable-binding level, language-particularly, without becoming crosslinguistic categories. The position is closer to Haspelmath's (2010, p. 666) reply to Newmeyer (formal references can be comparative concepts in their own right) than to the categorial-universalist alternative.

The metaphysical conclusion in §7, that the umbrella schema is a tool for locating projectibility profiles rather than itself a kind, is consistent with this. Three levels stay distinct: the comparative concept (mediated accessibility) lives at the cross-linguistic methodological level; descriptive categories (English number-transparent quantificational nouns, Korean case-stacking and scrambling, and so on) live at the language-particular grammatical level; PROJECTIBILITY PROFILES are observable bases licensing non-trivial projection, stabilised by identifiable sources, realised within or across such descriptive categories. Kindhood is a downstream verdict on profiles that earn it. The paper doesn't collapse the levels.

6.5 METHODOLOGICAL DEFAULT

Haspelmath's distinction has been contested in the typological literature; the present paper doesn't adjudicate. Haspelmath's position is taken as the safer methodological default for cross-linguistic claims, since it requires no commitment to category universality and doesn't assume a substantive crosslinguistic ontology of grammatical categories. Readers more comfortable with categorial universalism can read the umbrella schema as a useful working concept whose status is secondary to its methodological pay-off.

Croft's Radical Construction Grammar (RCG) is the construction-grammar variant closest to the present categorial-particularist commitment. RCG routes cross-linguistic comparison through language-particular constructions linked by similarity rather than through universal grammatical categories or formal apparatus. The umbrella schema's relation between comparative concept and language-particular descriptive category is broadly compatible with RCG's architecture: language-particular constructions instantiate the schema's variables (*X*, *P*, *R*) without committing the analyst to a universal category of transparency. The choice between routing through Haspelmath-style comparative concepts and routing through RCG's construction-similarity hierarchies is left as an option for further development; both preserve the categorial-particularist commitment.

6.6 BRIDGE TO §7

§7 develops the metaphysical claim: the umbrella schema locates projectibility profiles, each stabilised by an identifiable source (lexical convention, constructional inheritance, usage-based entrenchment, processing architecture, formal-semantic environment); Slater's (2015) SPC framework and Khalidi's (2018) causal-network account are diagnostic tools for assessing projection-supporting stability rather than master ontologies. The claim is consistent with categorial particularism: descriptive

categories are language-particular; the comparative concept *MEDIATED ACCESSIBILITY* lives at the cross-linguistic level; *PROJECTIBILITY PROFILES* are observable bases for projection, realised within or across descriptive categories. Cross-linguistic projection at the profile level isn't entailed by cross-linguistic projection at the comparative-concept level.

7 PROJECTIBILITY PROFILES

Section 1 set the paper's central question: when does a transparency attribution support induction, and what stabilises that induction? The previous sections worked through several candidates. Section 7 collects the answers. §7.1 explains why the umbrella schema doesn't itself name a single profile; it's a tool for locating profiles. §7.2 collects the projectibility profiles the paper has examined. §7.3 names the stabilising structures behind each. §7.4 places Slater's (2015) SPC framework and Khalidi's (2018) causal-network account as diagnostic tools rather than master ontologies. §7.5 states the resulting claim. §7.6 connects to related work and acknowledges limits. §7.7 bridges to §8.

7.1 THE UMBRELLA ISN'T A SINGLE PROFILE

Section 6 established that the umbrella schema is a comparative concept, an analyst's tool for asking the projectibility question of any language. The metaphysical question is what *follows* from posing the question across a heterogeneous family of cases. The umbrella's heterogeneity is the answer. The *X*-values across the family include lexical classes (number-transparent quantificational nouns), constructions (the ditransitive, partly-filled idioms), compounds (NN compounds), derivationally complex words (*-ness* and *-ost*' nominalisations), and intensional contexts (referential transparency). The *P*-values include number/count features, constituent meanings and their relations, compositional contributions, lexical-stem identity, and referents. The *R*-values include subject-verb agreement, interpretation, masked priming, and substitution under co-reference. No single stabilising source unifies these triples, and no single inductive base supports projection across them.

The umbrella thus identifies many candidate profiles, not one. The schema doesn't by itself license projection: knowing that a case instantiates the schema doesn't let us predict anything substantive about it, because the projectible content lives at the level of the individual profile, not the level of the schema. The schema's role is to make the projectibility question askable, case by case.

A useful general distinction follows from this heterogeneity. A *distributional class* names a real grouping at the surface — the schema's *X/P/R* slots get filled by some real configuration — while a *mechanism-maintained category* additionally requires some identifiable mechanism keeping the cluster projectible. The umbrella schema names a distributional class but not a mechanism-maintained category: the schema picks out a real family of cases, but no mechanism unifies them at the umbrella level. Each profile examined below has to be evaluated separately for whether it qualifies as a category in this stronger sense. The negative-polarity-item literature offers a parallel pattern: NPIs share a surface licensing property but the licensing patterns Hoeksema catalogues aren't manifestations of one mechanism, so NPIs are a class but not a category at the umbrella level (developed at length in (Reynolds, 2026b, ch. 7)).

7.2 TWO PROJECTIBILITY PROFILES

The paper has examined two projectibility profiles in detail.

NUMBER-TRANSPARENT QUANTIFICATIONAL NOUNS (§2.4). The observable base is class membership, identified by lexical identity, quantificational semantics, corpus distribution as head of an

of-complement construction, and typical determiner frame. The projected targets are non-trivial: near-obligatory agreement override in the central cases (in either direction); characteristic determiner patterns by subgroup; partitive vs. non-partitive *of*-complement availability; behaviour under fused-head construction; count/non-count percolation of the *of*-complement; restricted premodification. Class membership generates expectations for *new* members (boundary cases like *majority* and *minority*; gradient cases like *bunch*) that aren't built into the class's defining criteria. This is the strongest candidate profile examined here: the cluster supports inductive extension, not just cataloguing of the enumerated members.

The form-meaning accessibility family (§§3–5). Three within-domain profiles, related by analytical theme rather than by demonstrated cross-domain projection. (1) A construction-level profile: graded compositionality predicts compositional interpretability via inheritance and network position. (2) A compound-level profile: corpus-statistical predictors (N_1 frequency, semantic-relation expectedness given N_1 , N_2 productivity) predict rated transparency. (3) A derivationally-complex-word-level profile: rated transparency predicts long-SOA priming magnitude. The three differ in evidential strength: the compound profile rests on Bell and Schäfer's measured corpus predictors and rating data, the priming profile on Heyer and Kornishova's direct experimental evidence, while the construction-level profile is currently theoretical, with the network-transparency hypothesis offered as a hypothesis to be tested rather than as established empirical content. The three are compatible with a shared form-meaning accessibility hypothesis (§5.4). Whether the family constitutes a cross-domain projectibility cluster, on shared items, with cross-modal projection beyond what the individual profiles predict, is an open empirical question. On present evidence the family doesn't license such cross-domain projection. Appendix D sketches a within-items probe (corpus predictors, ratings, priming magnitude, and interpretation accuracy on the same set of derived nominalisations or compounds) whose outcome would settle the cross-domain question one way or the other. The form-meaning accessibility family is therefore the open research programme of the paper, not a delivered result.

Acquisition trajectory is a plausible further extension consistent with the network claim of §3.3 but not directly grounded by the present sections; sourcing usage-based acquisition evidence is left for further work.

7.3 STABILISING STRUCTURES

The two profiles are stabilised by different structures, and the paper's six *R*-values across §§2–5 (plus the boundary case of §4.3) call for a small inventory of stabilising sources rather than a single mechanism.

LEXICAL CONVENTION stabilises the number-transparent class qua class. Lexical inheritance here means that speakers acquire not isolated words but item-specific constructional slots: *a lot of X*, *plenty of X*, *the rest of X*, *a number of X*, with their associated agreement and complement-selection behaviour. The acquisition target is the slot, not just the lexeme. Speakers acquire a small, relatively stable inventory of these slots; new members enter at the boundary (*majority*, *minority*) or by sense-split (*couple₂*). Class identity itself isn't held by Boyd-style homeostatic feedback; it has the stability characteristic of any small, conventionally established grammatical class. The count/non-count percolation behaviour the class licenses connects to broader patterns in English countability; (Reynolds, 2026a) develops a Khalidian causal-cluster analysis of those patterns with conventional individuation as a central node. The transparency paper draws on that analysis but doesn't depend on its stronger claims; the within-NP percolation can be observed and predicted independently.

CONSTRUCTIONAL INHERITANCE stabilises the network-relative transparency hypothesis of §3.3. A construction’s compositional accessibility depends partly on the inheritance hierarchy it sits in; constructions whose inherited supercategories are well-attested support stronger compositional projection than orphaned constructions.

USAGE-BASED ENTRENCHMENT stabilises the corpus-statistical predictors of compound transparency (§4). Bell and Schäfer’s frequency, productivity, and semantic-relation expectedness predictors are corpus-derivable measurements of usage entrenchment; the rated-transparency target is a steady-state property whose stability is plausibly a reflex of that entrenchment.

PROCESSING ARCHITECTURE stabilises the time-course modulation of priming (§5). Heyer and Kornishova’s SOA-dependent transparency effect is a processing-architecture observation: the time-course of morphological decomposition makes morpho-orthographic information available before morpho-semantic information, and the resulting priming pattern depends on representational properties that the processing architecture exposes only at certain stages.

FORMAL-SEMANTIC ENVIRONMENT stabilises the referential transparency cases (§4.3). Substitution-salva-veritate behaviour depends on the logical structure of intensional contexts. This is the boundary case for the schema: the mediator is a context, not a form, and the stabilising source is the formal-semantic structure of the embedding.

The list isn’t closed. Other linguistic profiles may be stabilised by combinations of these or by sources the paper doesn’t invoke (information-structural pressures, learnability constraints, normative/pedagogical regularisation). The point is that “what stabilises this profile?” is the right empirical question, and the answer is heterogeneous.

7.4 SPC AND KHALIDI AS DIAGNOSTICS

SPC and Khalidi are complementary diagnostics rather than nested ones. Slater asks whether the cluster is stable and projection-supporting and is neutral on *what* stabilises it. Khalidi asks whether the stabilising structure is itself causal-hierarchical, with named candidate core properties whose causal action generates derivative properties. A profile can satisfy SPC without satisfying Khalidi, and (in principle) the reverse. The PROJECTIBILITY PROFILE is the unit of analysis; the two diagnostics ask different questions of it.

Slater (2015) gives the stability-plus-projection diagnostic. A profile satisfies SPC when its cluster of properties is stable enough to support reliable induction, regardless of *what* maintains the stability. Slater’s framework has affirmative content beyond maintenance-neutrality: cluster stability has to be cliquish (members co-occur reliably under counterfactually relevant conditions), and sub-cluster-to-cluster probabilistic entailment is what supports reliable projection. Number-transparent quantificational nouns are the strongest candidate to satisfy SPC examined here: the cluster is stable, the sub-group covariance documented in §2.4 (determiner-pattern, partitive availability, override-direction co-vary by sub-group) is the cliquishness SPC requires, and the maintenance is at least conventional/lexical for class identity (§7.3). The within-domain audit beyond override direction is empirically open; the cliquishness claim is licensed by the §2.4 sub-group covariance but not yet by an end-to-end corpus or experimental study. Each form-meaning window (§§3–5) satisfies SPC within its domain. Whether the family as a whole satisfies SPC at the cross-domain level depends on the empirical question §7.2 leaves open and §D sketches as a probe design.

Khalidi (2018) applies where the stabilising structure is ordered and causal-hierarchical. The account requires named candidate core properties whose causal action generates derivative properties

via identifiable pathways. The strongest sub-case in the paper isn't the QN class qua class but the count cluster the class participates in via *of*-complement percolation. Reynolds (2026a) analyses that cluster as Khalidian, with CONVENTIONAL INDIVIDUATION (speakers' construal of referents as discrete units, established by usage-based convention) as the candidate core property. The derivative properties causally linked to individuation include the count grammar's LOOSE features (plural agreement, the *many/few* contrast, plural-only forms) and TIGHT features (singular determiner *a(n)*, low-cardinal selection, count/non-count compatibility on determiners and predicates). Loose features can be satisfied by individuation alone; tight features require the precise atomic enumeration individuation supplies under stronger interpretation.

The number-transparent QN construction extends this chain at the construction level. The QN's *of*-complement determines the count/non-count profile of the whole NP (*a lot of money* non-count; *a lot of tourists* count), projecting the complement's individuation profile to the QN's external interface for agreement, partitivity, premodification, and predicate selection. Within-NP percolation is the causal-hierarchical move at the construction-internal level; the count cluster is the causal-hierarchical move at the lexical-grammatical level. Both have core/derivative structure in Khalidi's sense, and the construction's design isn't an analyst's add-on: it's a constructional pathway that routes the complement's individuation to the matrix relation, with predictable behavioural signatures (the cluster of §2.4) attached.

For the QN class qua class, Khalidi-style refinement isn't motivated on present evidence: class identity itself is conventionally stabilised without a separate ordered causal hierarchy. SPC is the diagnostic that fits class identity; Khalidi applies one level out, where the class plugs into the count cluster.

For the form-meaning family of §§3–5, Khalidi is acknowledged as a possible refinement *if* corpus predictors, rated transparency, priming magnitude, and interpretation accuracy align on shared items as joint reflexes of underlying representational structure. The four windows are not a temporal chain — masked priming probes early lexical access before explicit interpretation, and ratings are a late meta-linguistic judgement — but candidate joint consequences of a single underlying organisation. The within-items evidence doesn't yet exist; Appendix D sketches a probe design. Until that probe is run, Khalidi remains a possible refinement for the form-meaning family rather than a delivered analysis.

Kindhood, on this picture, is the downstream verdict. A profile that scores well on the SPC diagnostic, and especially one that also scores well on the Khalidi diagnostic, has earned the kind label. A profile that scores modestly is a candidate. A profile that fails either diagnostic isn't a kind, even if it remains a useful descriptive category. The diagnostic work is *what stabilises projection*; the kind label follows.

What does this diagnostic apparatus buy that plain projectibility-talk wouldn't? Three things. First, a maintenance-neutral test: Slater's SPC admits conventional, entrenched, architectural, and formal-semantic stabilisers on the same footing, where Boyd's homeostasis would force one mechanism (a poor fit for lexical convention) and an in-house neutrality criterion would have to be coined to play the same role. Second, a named diagnostic for causal-hierarchical structure: Khalidi gives the paper a way to ask whether the form-meaning family's sub-profiles have internal ordering, a question that would otherwise have to be raised without an existing vocabulary. Third, the class/category contrast of §7.1: the distinction between a distributional class (mere grouping) and a maintained-and-projectible category lets the paper distinguish profiles that name kinds from labels that merely

group co-occurring features. Where a profile's projectibility claim survives un-translated into the class/category vocabulary (number-transparent quantificational nouns are the clearest case), the diagnostic is light overhead. Where the claim leans on the diagnostic (the form-meaning family's possible causal-hierarchical structure), the apparatus earns its keep.

7.5 THE RESULTING CLAIM

Linguistic TRANSPARENCY isn't a kind, and the umbrella schema doesn't aspire to be one. The schema is a diagnostic tool for locating projectibility profiles: cases where an observable base licenses non-trivial projection to a target, stabilised by some identifiable source. Different transparency profiles are stabilised differently (lexical convention, constructional inheritance, usage entrenchment, processing architecture, formal-semantic environment), and each profile's projectibility is what makes the transparency attribution scientifically useful, not the profile's classification within any single ontology of kinds.

Three levels stay distinct, per §6:

- COMPARATIVE CONCEPT: mediated accessibility, the cross-linguistic methodological tool defined by the schema *X is transparent w.r.t. P for R when P remains accessible through X for the purposes of R*.
- DESCRIPTIVE CATEGORY: language-particular grammatical content. English number-transparent quantificational nouns are a CGEL category; English NN compounds and *-ness* nominalisations are descriptive categories defined by language-particular criteria; Korean case-stacking and scrambling are Korean-particular categories; and so on.
- *Projectibility profile*: an observable base that licenses non-trivial projection to a target, stabilised by some identifiable source, realised within or across language-particular descriptive categories. This is the *category* of §7.1: a maintained-and-projectible cluster, contrasted with a *class* (mere distributional grouping). Profiles that score well on the SPC and (where appropriate) Khalidi diagnostics earn the local-kind label as a downstream verdict.

The contribution here is to apply the projectibility-first move to a *relational* concept. TRANSPARENCY is the relation *X-mediates-P-for-R*, not an entity. Asking *when does this attribution support induction, and what stabilises the induction?* works for relations as it does for the umbrella concepts *language* and *grammatical category* addressed in related work. Kindhood enters the picture as a verdict on profiles that earn it, not as the master question that organises the analysis.

7.6 CONNECTIONS AND LIMITS

Other multi-sense umbrella concepts (e.g. *language*, *grammatical category*) admit the same projectibility-first move: the umbrella is not itself a kind but specific senses can be projectible profiles. Reynolds (2026b, ch. 7) develops the structural analogue for negative-polarity items at length.

What would falsify the projectibility profiles the paper has examined? For number-transparent quantificational nouns: if the predicted behaviours (override, determiner restrictions, partitive availability, premodification) failed to track class membership in a substantial corpus or experimental study, the SPC diagnostic would fail. For the form-meaning family: if rated transparency, priming magnitude, and interpretation accuracy failed to align as joint consequences of underlying representational structure on the same items in a single study, the cross-domain projection question would resolve negatively, and the family would remain three within-domain profiles rather than becoming

a single cross-domain profile. The empirical bar is real, and the within-domain checks of §§2.4, 4.4, and 5.5 are first passes rather than final verdicts.

The Korean (§2.6) extension is the most consequential gap. If the comparative concept of mediated accessibility fails to yield productive cross-linguistic generalisation when applied to Korean case-stacking and scrambling, the paper's typological discipline (§6) does what discipline is supposed to do: it admits the negative result without sacrificing the comparative-concept frame. The methodological discipline survives even if particular case studies don't extend cleanly.

7.7 BRIDGE TO §8

§8 closes by returning to the question of §1. The cases the paper worked through were tested as projectibility profiles, not classified as kinds. Some are strong candidates (number-transparent NPs, on the strength of the §2.4 projection set). Some are within-domain profiles whose cross-domain integration is an open research programme (the form-meaning family). Some are non-instances (coordination) or boundary cases (referential transparency). The paper's claim is that TRANSPARENCY is a disciplined comparative concept whose attributions are scientifically useful when they identify a projectibility profile and locate the source of its stability.

8 CONCLUSION

Linguistic transparency is best understood as a family of mediated accessibility relations rather than as a single syntactic property. The umbrella schema, *X is transparent with respect to property P for relation R when P remains accessible through X for the purposes of R*, is a comparative concept: an analyst's question, applicable to any language, that asks which configurations let an internal property remain available to a grammatical or interpretive relation. The schema doesn't claim that any language has a descriptive category called TRANSPARENCY. It states the question.

The conclusion developed in §7 is that the umbrella schema locates PROJECTIBILITY PROFILES rather than naming a single kind. The *X/P/R* values across the family are heterogeneous: no shared stabilising source unifies them, and the projectible content lives at the level of the individual profile, not the level of the schema. The schema's job is to make the projectibility question askable; profiles earn scientific use when they identify an observable base licensing non-trivial projection to a target, stabilised by an identifiable source.

The profiles the paper has worked through:

- NUMBER-TRANSPARENT QUANTIFICATIONAL NOUNS (§2.4). A small, relatively stable category in the descriptive system of CGEL English (qua linguistic class, qua maintained-and-projectible category in §7.1's stronger sense, once maintenance is attributed). Class membership predicts agreement override, characteristic determiner patterns, partitive-vs-non-partitive *of*-complement availability, behaviour under fused-head construction, count/non-count percolation, and premodification tendencies. Stabilised by conventional lexical inheritance. The strongest candidate to pass Slater's (2015) broad stability-plus-projection diagnostic among the cases examined; the within-domain audit is open.
- *The form-meaning accessibility family* (§§3–5). Three within-domain projectibility profiles related by analytical theme rather than by demonstrated cross-domain projection: a construction-level profile stabilised by constructional inheritance (CxG); a compound-level profile stabilised by usage-based entrenchment (Bell and Schäfer); a derivationally-complex-word-level profile stabilised by processing architecture (Heyer and Kornishova). Each licenses within-domain

projection. Whether the three compose into a cross-domain projectibility cluster on shared items, and whether Khalidi (2018)-style causal-network refinement applies to that cluster, depend on within-items cross-modal evidence the present paper doesn't provide.

Coordination is excluded as a non-instance (§2.3): nothing was going to be encapsulated. Referential transparency is acknowledged as the boundary case (§4.3) where the schema's mediator is a context rather than a form; the philosophical-semantics tradition is named without being developed.

What is left open. Korean case-stacking and scrambling (§2.6) are deferred for separate co-authored development; the comparative-concept frame licenses asking the schema's question of those patterns with *R*-values different from agreement (argument-role recovery, scope, information-structure assignment), but the language-particular content is pending. The within-items cross-modal evidence that would test whether the form-meaning family has Khalidi-style causal-hierarchical structure (corpus predictors, ratings, priming at multiple SOAs, and interpretation accuracy on a single item set) is future work. Acquisition trajectory is a plausible further extension consistent with usage-based theories but not directly grounded by the present cases.

The paper's closing claim. TRANSPARENCY in linguistics names a family of mediated accessibility relations that the umbrella schema locates as candidate projectibility profiles. A transparency attribution is scientifically useful when it identifies a profile: an observable base that licenses non-trivial projection to a target, stabilised by an identifiable source. Some profiles support strong projection; others weak; some none. Slater's SPC framework and Khalidi's causal-network account are diagnostic tools for assessing how well a profile supports induction. Kindhood is the downstream verdict on profiles that earn it, not the master question of the analysis.

A NUMBER-TRANSPARENT NP EMPIRICAL APPARATUS

This appendix collects the COCA pilot data behind the §2 claims about the number-transparent quantificational noun class. The body of the paper carries the qualitative claims; the audit trail lives here. Full machine-readable data: `paper/data/coca-pilot.md`.

A.1 METHODOLOGY

Queries were issued through the local Playwright wrapper at `tools/english-corpora/bin/ecorg.mjs` (May 2026). Each query is a surface string; counts are raw COCA frequencies, not filtered for head versus modifier use, idiom contamination, or sense disambiguation, except where noted. Two filtering issues recur:

- *Modifier confound*. Surface strings of the form *X of N* may include cases where *N* is a modifier inside a larger NP rather than the head of the *of*-complement (e.g., *lots of day-care* rather than *lots of day*). The countability paper (Reynolds, 2026a) handles the same coding issue for *three police* and *three cattle*. Where small counts fall in the predicted-counter direction, the modifier confound likely inflates them.
- *Parse-shift*. Surface strings containing a number-transparent QN plus a verb form may include cases where the QN sits inside a relative clause modifying a different head-noun, with the visible verb agreeing with that other head. *A lot of money are* is the textbook case: of 13 raw plural-agreement hits, 10 are parse-shifted (people/parents/movies/times *are*; *a lot of money* in a relative clause), 1 is irrealis *were* (CGEL §3.5.1), 1 3 hit is contracted *has* (singular, matches prediction), and 1 is a genuine plural override.

The wrapper's KWIC mode (which would let an analyst inspect contexts) is broken as of May 2026 (Playwright internal error on every attempt). Counts in the tables below were generated via list-mode queries; the parse-shift filtering for *a lot of money* was done from contexts retrieved separately. Other small-count counter-direction cells in the override table have not been spot-checked and may include similar artefacts.

A.2 A.2 PARTITIVE VS. NON-PARTITIVE OF-COMPLEMENTS

The contrast: PARTITIVE *of*-complements have definite complements (*of the money*, *of those books*); NON-PARTITIVE *of*-complements have indefinite complements (*of money*, *of soldiers*).

A.2.1 Plural-only forms (*lots*, *bags*, *heaps*, *loads*, *oodles*, *stacks*, *piles*)

Non-partitive (X of N): *lots of people* 3,878; *lots of money* 1,376; *lots of time* 703; *loads of money* 109; *piles of money* 91; *loads of people* 90; *loads of time* 34; *heaps of money* 15; *heaps of people* 7; plus single-digit hits for several other combinations. Total: 6,311. The single-digit *X of day* hits (*lots of day* 3, *heaps of time* 3, *piles of time* 1, *piles of people* 1) likely include modifier uses (*lots of day-care*).

Partitive (X of the N): *lots of the time* 7; *lots of the people* 5; *loads of the day* 1. Total: 13. *Piles* and *heaps* return 0 partitive hits with any complement tested.

Indefinite-NP complement (X of a N): *lots of a day* 0; *piles of a day* 0; *heaps of a day* 0; *loads of a day* 0; *lots/piles/heaps/loads of an afternoon* 0. Total: 0. Plural-only forms reject *of [Det NP]* whether the determiner is definite (*the*) or indefinite (*a/an*); their *of*-complements are bare/non-determined.

Partitive ratio: $13 / (6,311 + 13) \approx 0.2\%$.

A.2.2 *Plenty*

Non-partitive: *plenty of time* 4,017; *plenty of people* 1,563; *plenty of money* 668; *plenty of day* 4 (likely modifier). Total: 6,252.

Partitive: *plenty of the time* 1; *plenty of the people* 1. Total: 2.

Partitive ratio: $2 / 6,254 \approx 0.03\%$.

A.2.3 *Rest, remainder*

Non-partitive (the X of N): *the rest of time* 75; *the rest of people* 12; *the rest of day* 7; *the remainder of time* 3; *the remainder of day* 2; *the rest of money* 1. Total: 100 raw. *The rest of time* (75 hits) is heavily contaminated by idiomatic *for the rest of time* (= “forever”) and similar. Spot-checking head uses only would lower the figure substantially.

Partitive (the X of the N): *the rest of the day* 1,932; *the rest of the time* 830; *the rest of the people* 272; *the rest of the money* 222; *the remainder of the day* 75; *the remainder of the time* 21; *the remainder of the money* 13. Total: 3,365.

Indefinite-NP complement: *the remainder of a day* 1.

Partitive (definite NP) ratio: $3,365 / (3,365 + 100 + 1) \approx 97.1\%$ raw, higher after filtering the *for the rest of time* idiom.

A.3 OVERRIDE-DIRECTION AGREEMENT

Test of the “near-obligatory in central cases” claim. Subject NP plus finite verb agreement, by complement number/count.

Subject	Predicted	Plural agr.	Singular agr.	% predicted
<i>a lot of people</i>	plural	4,195 (ARE 3,406; WERE 789)	85 (IS 72; WAS 13)	98.0%
<i>a lot of money</i>	singular	1 (ARE, filtered)	90 (IS 61; WAS 29)	98.9%
<i>a number of people</i>	plural	100 (ARE 52; WERE 48)	4 (IS 2; WAS 2)	96.2%
<i>lots of people</i>	plural	348 (ARE 288; WERE 60)	6 (IS 5; WAS 1)	98.3%
<i>lots of money</i>	singular	1 (ARE 1)	24 (IS 19; WAS 5)	96.0%
<i>plenty of people</i>	plural	79 (ARE 65; WERE 14)	0	100%
<i>plenty of money</i>	singular	0	6 (IS 4; WAS 2)	100%
<i>the rest of the people</i>	plural	19 (ARE 15; WERE 4)	0	100%
<i>the rest of the money</i>	singular	0	19 (IS 10; WAS 9)	100%

Override is in the predicted direction at 96–100% across all tested QN+complement combinations after filtering. *A lot of money* required filtering: of 13 raw plural-agreement hits, 10 were parse-shifted (the surface string *a lot of money* are appearing inside relative clauses modifying a plural-noun head, e.g., *people who earn a lot of money are successful*; *movies that make a lot of money are the biggest help*; *times when we... make a lot of money are gone*), 1 was the irrealis *were* (*if a lot of money were at stake*), and 1 's hit was singular *has* (matches prediction); only 1 hit was a genuine plural override (*a lot of money are being raised*). Filtered count: 1 plural / 90 singular = 98.9% singular.

Other small-count counter-direction cells (*a number of people is/was* 4 hits; *lots of people is/was* 6 hits; *lots of money are* 1 hit) have not been spot-checked and may include similar artefacts; the headline pattern doesn't depend on filtering them.

A.4 A.4 BUNCH: ANIMATE VS. INANIMATE AGGREGATES

CGEL claims that *bunch* tilts collective for inanimate aggregates and transparent for groups of people. The COCA pilot supports the animate half but not the inanimate half.

Subject	Total tokens	<i>was</i>	<i>were</i>
Animate aggregates			
<i>a bunch of people</i>	1,421	0	28
<i>a bunch of kids</i>	448	0	12
<i>a bunch of hooligans</i>	12	0	0
Inanimate aggregates			
<i>a bunch of flowers</i>	68	0	0
<i>a bunch of cars</i>	–	0	1
<i>a bunch of papers</i>	–	0	0

Subject	Total tokens	<i>was</i>	<i>were</i>
<i>a bunch of leaves</i>	–	0	0
<i>a bunch of things</i>	–	0	1

Animates: 40 plural / 0 singular = 100% plural. The CGEL “transparent for groups of people” claim is supported.

Inanimates: 2 plural / 0 singular. The CGEL “tilts collective for inanimate aggregates” claim is not supported in COCA. Inanimate-bunch is mostly absent from subject position (typical use is as object: *she gave him a bunch of flowers*); the two inanimate-bunch hits in subject position both took plural rather than the CGEL-predicted singular. *A bunch of flowers was presented to the teacher* (the CGEL example) returns 0 hits in COCA. KWIC inspection of the 40 animate-plural hits for parse-shift is pending the wrapper fix.

B COMPOUND TRANSPARENCY: BELL AND SCHÄFER'S EMPIRICAL APPARATUS

This appendix carries the detailed exposition of Bell and Schäfer (2016) that §4 references in compressed form. The body keeps the headline finding (three predictors, scalar transparency, within-domain projection); the predicate inventory and stress speculation live here.

B.1 LEVI'S PREDICATE INVENTORY

Bell and Schäfer code the semantic relation between the constituents of each NN compound using Levi's (1978) nine predicates: CAUSE, HAVE, MAKE, USE, BE, IN, FOR, FROM, ABOUT (Bell & Schäfer, 2016, p. 167).

Three of the predicates (CAUSE, HAVE, MAKE) come in two variants depending on which constituent functions as the subject of the predicate:

- CAUSE₁ *tear gas* (“gas that causes tears”) versus CAUSE₂ *drug deaths* (“deaths caused by drugs”).
- HAVE₁ *picture book* versus HAVE₂ *government land*.
- MAKE₁ *honeybee* versus MAKE₂ *daisy chains*.

The other six predicates (USE, BE, IN, FOR, FROM, ABOUT) don't subdivide.

B.2 METAPHORIC/METONYMIC SHIFT CODING

Bell and Schäfer's earlier work coded each compound and each constituent for metaphoric and metonymic shifts, treating shift as a binary feature (Bell & Schäfer, 2016, p. 167). The coding is independent of the predicate-relation assignment and contributes separately to predicting transparency ratings.

B.3 IN, FOR, AND THE STRESS HYPOTHESIS

Two of the Levi relations are positively associated with whole-compound transparency in Bell and Schäfer's models: IN and FOR. Drawing on Bell and Plag (2012), Bell and Schäfer suggest a possible mechanism: if more transparent constituents are less informative, and if compound stress falls on the more informative constituent, the IN/FOR effects on transparency would line up with the stress patterns observed for those compounds (Bell & Schäfer, 2016, p. 168). The connection is offered as a hypothesis that further work would need to test, not as an established mechanism.

B.4 N₁/N₂ ASYMMETRY

Although whole-compound transparency is a function of the transparencies of the two constituents, N₁ (the modifier) and N₂ (the head) differ in the nature of their contribution. The asymmetry is one of the two headline conclusions of Bell and Schäfer's paper. The other is that all the significant predictors of rated transparency in their models are also independent predictors of processing speed, leading them to a tentative hypothesis: "perceived transparency may itself be a reflex of ease of processing" (Bell & Schäfer, 2016, 157, abstract). The hypothesis is offered as a direction for further work, not a result of the present study.

B.5 THE REDDY ET AL. DATABASE

The empirical anchor of Bell and Schäfer's models is the database of Reddy and colleagues (2011), used as the dependent variable. The database consists of 30 transparency ratings for each of 90 two-part English NN compounds, totalling 8,100 ratings, collected via Amazon Mechanical Turk. Annotators rated, on a 0–5 scale, how literal each constituent was in the compound and how literal the compound as a whole was. Ratings are continuous; they don't partition compounds into transparent-versus-opaque classes (Bell & Schäfer, 2016, p. 164).

C MORPHOLOGICAL PRIMING: HEYER AND KORNISHOVA'S DESIGN AND STATISTICAL DETAIL

This appendix carries the experimental detail from Heyer and Kornishova (2018) that §5 references in compressed form. The body keeps the design logic (short vs. long SOA; transparency continuous; key result that transparency modulates priming only at long SOAs); participant counts, exact SOAs, item-selection, and statistical results are here.

C.1 PARTICIPANTS AND DESIGN

Two masked-priming experiments, one with English *-ness* nominalisations (49 native speakers: 24 in the short-SOA condition, 25 in the long-SOA condition) and one with Russian *-ost'* nominalisations (61 native speakers: 31 short-SOA, 30 long-SOA) (Heyer & Kornishova, 2018, p. 1115).

Targets in both experiments were the adjective bases (e.g., *pale*, *dark*). Primes came in three types:

- the corresponding nominalisation as a related prime (e.g., *paleness* preceding *pale*);
- a simplex noun as an unrelated control prime;
- the adjective itself as an identity prime.

Each experiment used two stimulus-onset asynchronies (SOAs):

- short: 39 ms English, 33 ms Russian;
- long: 77 ms English, 67 ms Russian.

C.2 ITEMS AND TRANSPARENCY RATING

Items spanned the full transparency range. At the opaque end, *business* (rated 3.17) is no longer interpretable as "state of being busy"; at the transparent end, *falsehood* (rated 5.75) is fully derivable from *false*; intermediate items occupy the middle of the scale.

Out of 95 *-ness* items pre-tested, 30 were selected for the main experiment to cover the range continuously (Heyer & Kornishova, 2018, p. 1115). The Russian experiment used a parallel design with 30 *-ost'* nominalisations spanning a similar range.

Transparency was rated by separate groups of native speakers (28 English, 38 Russian) on a 1–7 Likert scale for semantic relatedness.

The two suffixes are highly comparable: both deadjectival, both producing abstract nouns denoting “state of being *X*” (e.g., *dark* + *ness* = *darkness*; *bledn-* + *-ost'* = *blednost'* “paleness”), and both productive (Heyer & Kornishova, 2018, p. 1115).

C.3 STATISTICAL RESULTS

At short SOA in both languages, the data show constant morphological priming across the transparency scale. The interaction between Prime Type and Transparency doesn't reach significance: priming effects are 37 ms in English and 16 ms in Russian, but they aren't modulated by how transparent the prime-target pair is rated (Heyer & Kornishova, 2018, pp. 1117–1119).

At long SOA, priming becomes transparency-modulated: more transparent prime-target pairs (e.g., *paleness-pale*) yield larger priming effects than less transparent pairs (e.g., *business-busy*). The three-way interaction between Prime Type, Transparency, and SOA is significant in English ($t = 2.10$, $p = 0.036$) and marginally so in Russian ($t = -1.69$, $p = 0.091$; the negative sign reflects factor coding rather than a reversed effect, and the direction of transparency-on-priming is consistent across the two languages). The English result establishes the SOA-dependent modulation; the Russian result trends in the same direction without reaching conventional significance (Heyer & Kornishova, 2018, pp. 1117–1119).

C.4 RIVAL PARALLEL-PROCESSING POSITION

The serial reading is contested. Feldman et al.'s (2015) *Must analysis of meaning follow analysis of form?* is the canonical parallel-immediate paper: a time-course analysis with a same-target design (each target compared after semantically similar and semantically dissimilar morphologically structured primes within subjects) and SOA treated as continuous from 34 to 100 ms. Their finding is that semantic similarity affects recognition latencies even at the shortest SOAs (34 and 48 ms), with the effect of shared semantics increasing linearly with SOA when long SOAs are intermixed. The serial form-then-meaning ordering predicts no semantic effect at SOAs in the 34–48 ms range; the parallel-immediate reading takes the observed effect as direct evidence against that prediction. Heyer and Kornishova's design (separate item sets at each SOA, transparency as a continuous predictor, two languages) gives the serial reading a different test: the absence of transparency modulation across the item set at the short SOA. The two designs probe different patterns. Same-target wedge (Feldman et al.) tests whether semantic similarity affects latencies on identical targets at very short SOAs; across-item gradient (Heyer and Kornishova) tests whether transparency rating predicts priming magnitude on a continuum. They aren't head-to-head, and the literature is mixed enough (the meta-table on p. 1121 of Heyer and Kornishova surveys 35 prior masked-priming studies) that the present paper takes the serial reading as Heyer and Kornishova's reading, not as a settled disciplinary verdict.

D A CROSS-MODAL PROBE DESIGN FOR THE FORM-MEANING FAMILY

This appendix sketches the cross-modal probe that §§7.2 and 7.4 flag as the open empirical question for the form-meaning family. The aim is not a finished study but a tractable design whose outcome would settle the cross-domain projectibility question one way or the other.

D.1 HYPOTHESIS AND FALSIFICATION

The form-meaning family of §§3–5 comprises three within-domain profiles: theoretical compositionality, rated compound transparency, time-sensitive masked priming. Each licenses within-domain

projection. The cross-domain hypothesis is that the three are joint reflexes of a single underlying organisation — shared representational structure that exposes form-meaning composability through different windows. The hypothesis is non-trivial because alternative organisations are available: each profile could be domain-internal, with no cross-domain alignment. The probe asks whether the four measurement windows align on the *same items*, which is the only test that distinguishes the cluster reading from the three-separate-profiles reading.

Falsification is built in. If corpus predictors, rated transparency, priming magnitude, and interpretation accuracy fail to align on shared items as joint consequences of representational structure, the family is three within-domain profiles and the cross-domain projectibility question resolves negatively. The diagnostic is a single, item-anchored study; sample-size and statistical-power planning would follow the four-window design rather than any single-window precedent.

D.2 ITEM SET

Deadjectival nominalisations on *-ness* are the natural anchor: §5 establishes a within-domain transparency continuum with rated values from *business* (3.17) to *falseness* (5.75); the *-ness* items have been used in masked-priming work and are corpus-tractable; and parallel Russian *-ost'* items extend the design cross-linguistically without re-doing the architecture. A target set of 60–80 items spanning the rating continuum, with matched controls (frequency-matched simplex nouns, frequency-matched derivations on a different suffix), would support the four windows.

NN compounds (the §4 target) are an alternative or supplementary anchor. The same within-items logic applies, with Bell and Schäfer's (2016) corpus predictors providing the corpus measurement directly. Whether to run the probe on *-ness*, on NN compounds, or on both is a research-design choice; the probe's logical structure is identical across item types.

D.3 MEASUREMENTS

Four within-items measurements are required.

- CORPUS PREDICTORS. Frequency, productivity, and semantic-relation expectedness, computed per item from a balanced corpus (BNC, COCA, or COHA depending on the item type). Bell and Schäfer's measurement procedure for compounds extends naturally to nominalisations.
- RATED TRANSPARENCY. Native-speaker ratings on a 1–7 Likert scale, replicating Heyer and Kornishova's rating procedure with a separate participant group from the priming sub-study.
- PRIMING MAGNITUDE. Masked-priming reaction times at two SOAs (short, ≈ 30 –40 ms; long, ≈ 60 –80 ms), with the morphologically related prime, an unrelated simplex control, and an identity prime, replicating Heyer and Kornishova's design within the new item set.
- INTERPRETATION ACCURACY. A paraphrase or sentence-completion task probing whether speakers can recover the relation between the derivative and its base in context. The task should be designed so that opaque items predict reduced accuracy and longer response times relative to transparent items.

The four windows probe representational structure at distinct stages of processing and explicit access. Their alignment, if found, is a non-trivial empirical pattern.

D.4 STATISTICAL MODEL

A joint analysis with item-level random effects and a latent representational-structure factor is the natural statistical move. Each measurement is a noisy reflex; the latent factor's loadings index how

strongly each window reflects the underlying structure. The cluster hypothesis predicts non-trivial loadings on all four windows; the three-separate-profiles alternative predicts that the loadings partition by domain, with no shared factor.

The temporal asymmetry §7.4 flags between priming and interpretation is structural: priming probes early lexical access; interpretation probes explicit recovery; ratings are a metalinguistic judgement supplied at a separate, untimed task. The four windows aren't a temporal chain. The latent-factor model treats them as joint reflexes, with no claim that any one window causes any other.

D.5 STATUS

The probe doesn't exist. The form-meaning family is currently three within-domain profiles plus a research hypothesis. §7.4's Khalidi-refinement reading is conditional on the probe's outcome; absent the probe, Slater's SPC diagnostic is the relevant standard, and each within-domain profile is treated separately.

ACKNOWLEDGEMENTS

This paper was drafted with the assistance of large language models (Claude Opus 4.7 with the iM-context configuration; OpenAI Codex 5.3). All content has been reviewed and revised by the authors, who take full responsibility for the final text.

REFERENCES

- Bell, M. J., & Schäfer, M. (2016). Modelling semantic transparency. *Morphology*, 26(2), 157–199. <https://doi.org/10.1007/s11525-016-9286-3>
- da Cunha, Y., & Gibson, E. (2025). Syntactic complexity phenomena are better explained without empty elements mediating long-distance dependencies. *Cognitive Science*, 49(7), e70088. <https://doi.org/10.1111/cogs.70088>
- Feldman, L. B., Milin, P., Cho, K. W., Moscoso del Prado Martín, F., & O'Connor, P. A. (2015). Must analysis of meaning follow analysis of form? A time course analysis. *Frontiers in Human Neuroscience*, 9, 111. <https://doi.org/10.3389/fnhum.2015.00111>
- Gibson, E. (2025). *Syntax: A cognitive approach*. MIT Press.
- Goldberg, A. E. (2006). *Constructions at work: The nature of generalization in language*. Oxford University Press. <https://doi.org/10.1093/acprof:oso/9780199268511.001.0001>
- Grosu, A. (2003). A unified theory of standard and transparent free relatives. *Natural Language and Linguistic Theory*, 21(2), 247–331. <https://doi.org/10.1023/A:1023387825622>
- Grosu, A. (2016). The semantics, syntax, and morphology of transparent free relatives revisited: A comparison of two approaches. *Natural Language and Linguistic Theory*, 34(4), 1245–1280. <https://doi.org/10.1007/s11049-015-9323-7>
- Haspelmath, M. (2010). Comparative concepts and descriptive categories in crosslinguistic studies. *Language*, 86(3), 663–687. <https://www.jstor.org/stable/40961695>
- Heyer, V., & Kornishova, D. (2018). Semantic transparency affects morphological priming ... eventually. *Quarterly Journal of Experimental Psychology*, 71(5), 1112–1124. <https://doi.org/10.1080/17470218.2017.1310915>
- Huddleston, R., & Pullum, G. K. (2002). *The Cambridge grammar of the English language*. Cambridge University Press.

- Khalidi, M. A. (2018). Natural kinds as nodes in causal networks. *Synthese*, 195(4), 1379–1396. <https://doi.org/10.1007/s11229-015-0841-y>
- Kim, J.-B. (2026). *Form and function mismatches in English syntax: A construction-grammar perspective* (1st ed.). Routledge.
- Newmeyer, F. J. (2007). Linguistic typology requires crosslinguistic formal categories. *Linguistic Typology*, 11(1), 133–157. <https://doi.org/10.1515/LINGTY.2007.012>
- Quine, W. V. (1956). Quantifiers and propositional attitudes. *The Journal of Philosophy*, 53(5), 177–187. <https://www.jstor.org/stable/2022451>
- Reynolds, B. (2026a). *The Homeostatic maintenance of English countability: Bidirectional inference and the stability of grammatical clusters* [Manuscript submitted to *Linguistics*].
- Reynolds, B. (2026b). *Words that won't hold still: How linguistic categories work* [Manuscript in preparation]. <https://github.com/BrettRey/hpc-book>
- Reynolds, B. (2026c, March). *Two kinds of transparency: Category flexibility and the verb inventory in TFRs* [Working note. Empirical core to be developed in joint work with Jong-Bok Kim, in preparation.].
- Schelfhout, C., Coppen, P.-A., & Oostdijk, N. (2003). Transparent free relatives. In S. Blaho, L. Vicente & M. de Vos (Eds.), *Proceedings of the twelfth conference of the student organization of linguistics in europe (CONSOLE XII)*.
- Slater, M. H. (2015). Natural kindness. *British Journal for the Philosophy of Science*, 66(2), 375–411. <https://doi.org/10.1093/bjps/axto33>
- van Riemsdijk, H. (2006). Free relatives. In M. Everaert & H. van Riemsdijk (Eds.), *The Blackwell companion to syntax* (pp. 338–382). Blackwell.
- Wilder, C. (1999). Transparent free relatives. In K. N. Shahin, S. Blake & E.-S. Kim (Eds.), *Proceedings of the west coast conference on formal linguistics* (pp. 685–689, Vol. 17). CSLI Publications.